

(Leveled) Fully Homomorphic Encryption without Bootstrapping

ZVIKA BRAKERSKI, Weizmann Institute of Science
CRAIG GENTRY, IBM Research
VINOD VAIKUNTANATHAN, MIT and University of Toronto

We present a novel approach to fully homomorphic encryption (FHE) that dramatically improves performance and bases security on weaker assumptions. A central conceptual contribution in our work is a new way of constructing leveled, fully homomorphic encryption schemes (capable of evaluating arbitrary polynomial-size circuits of a-priori bounded depth), without Gentry's bootstrapping procedure.

Specifically, we offer a choice of FHE schemes based on the learning with error (LWE) or Ring LWE (RLWE) problems that have 2^λ security against known attacks. We construct the following.

- (1) A leveled FHE scheme that can evaluate depth- L arithmetic circuits (composed of fan-in 2 gates) using $\tilde{O}(\lambda \cdot L^3)$ per-gate computation, quasilinear in the security parameter. Security is based on RLWE for an approximation factor exponential in L . This construction does not use the bootstrapping procedure.
- (2) A leveled FHE scheme that can evaluate depth- L arithmetic circuits (composed of fan-in 2 gates) using $\tilde{O}(\lambda^2)$ per-gate computation, which is independent of L . Security is based on RLWE for quasipolynomial factors. This construction uses bootstrapping as an optimization.

We obtain similar results for LWE, but with worse performance. All previous (leveled) FHE schemes required a per-gate computation of $\tilde{\Omega}(\lambda^{3.5})$, and all of them relied on subexponential hardness assumptions.

We introduce a number of further optimizations to our scheme based on the Ring LWE assumption. As an example, for circuits of large width (e.g., where a constant fraction of levels have width $\Omega(\lambda)$), we can reduce the per-gate computation of the bootstrapped version to $O(\lambda)$, independent of L , by batching the bootstrapping operation. At the core of our construction is a new approach for managing the noise in lattice-based ciphertexts, significantly extending the techniques of Brakerski and Vaikuntanathan [2011b].

Categories and Subject Descriptors: E.3 [**Data**]: Data Encryption—*Public key cryptosystems*

General Terms: Security, Theory

Additional Key Words and Phrases: Fully homomorphic encryption, lattices, learning with errors

ACM Reference Format:

Brakerski, Z., Gentry, C., and Vaikuntanathan, V. 2014. (Leveled) fully homomorphic encryption without bootstrapping. *ACM Trans. Comput. Theory* 6, 3, Article 13 (July 2014), 36 pages.

DOI:<http://dx.doi.org/10.1145/2633600>

13

Z. Brakerski was supported by a Simons Postdoctoral Fellowship. C. Gentry was supported by DARPA under agreement number FA8750-11-C-0096. V. Vaikuntanathan was supported by an NSERC Discovery Grant and by DARPA under Agreement number FA8750-11-2-0225.

The U.S. Government is authorized to reproduce and distribute reprints for Governmental purposes notwithstanding any copyright notation thereon. The views and conclusions contained herein are those of the authors and should not be interpreted as necessarily representing the official policies or endorsements, either expressed or implied, of DARPA, or the U.S. Government. Approved for Public Release, Distribution Unlimited.

Authors' addresses: Z. Brakerski (corresponding author), Weizmann Institute of Science, Rehovot, Israel; email: zvik.brk@gmail.com; C. Gentry, IBM Research T.J. Watson, Yorktown Heights, NY; V. Vaikuntanathan, MIT Computer Science and Artificial Intelligence Laboratory, Cambridge, MA.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2014 ACM 1942-3454/2014/07-ART13 \$15.00

DOI:<http://dx.doi.org/10.1145/2633600>

1. INTRODUCTION

Fully homomorphic encryption (FHE) [Gentry 2009b; Rivest et al. 1978] allows a computationally powerful worker to receive encrypted data and perform arbitrarily complex, dynamically chosen computations on that data while it remains encrypted, despite not having the secret decryption key. Until recently, all FHE schemes [Brakerski and Vaikuntanathan 2011a; Coron et al. 2011; Gentry 2009b; Gentry and Halevi 2011b; Smart and Vercauteren 2010; van Dijk et al. 2010] followed the same blueprint, namely, the one laid out in Gentry’s original construction [Gentry 2009a, 2009b].

The first step in Gentry’s blueprint is to construct a *somewhat homomorphic encryption (SWHE) scheme*, namely, an encryption scheme capable of evaluating low-degree multivariate polynomials homomorphically. Starting with Gentry’s original construction based on ideal lattices [Gentry 2009b], there are by now a number of such schemes in the literature [Brakerski and Vaikuntanathan 2011a; Coron et al. 2011; Gentry and Halevi 2011b; Lauter et al. 2011; Smart and Vercauteren 2010; van Dijk et al. 2010], all of which are based on lattices (either directly or implicitly). The ciphertexts in all these schemes are “noisy”, where the noise grows slightly during homomorphic addition and explosively during homomorphic multiplication and, hence, the limitation of low-degree polynomials.

To obtain FHE, Gentry provided a remarkable *bootstrapping theorem* which states that given an SWHE scheme that can evaluate its own decryption function (plus an additional operation), one can transform it into a “leveled” FHE scheme. Bootstrapping “refreshes” a ciphertext by running the decryption function on it homomorphically using an encrypted secret key (given in the public key), resulting in a reduced noise. A leveled FHE scheme is a slightly weaker, but still practically as useful, variant of FHE, where the parameters of the scheme may depend on the *depth* of the circuits that the scheme can evaluate (but not on their size). The schemes we construct in this work are all leveled FHE schemes. One can obtain a “pure” FHE scheme (with a constant-size public key) from these leveled FHE schemes by assuming *circular security*, namely, that it is “safe” to encrypt the leveled FHE secret key under its own public key. With this understanding, and when there is no cause for confusion, we will omit the term “leveled” throughout this work.

Thus, to finish the construction, it is sufficient to design an SWHE scheme that is capable of homomorphically evaluating its own decryption circuit (plus some). Unfortunately, until very recently, natural SWHE schemes were incapable of evaluating their own decryption circuits without making significant modifications to the scheme, resulting in additional cryptographic assumptions. (We discuss two recent and important exceptions [Brakerski and Vaikuntanathan 2011a; Gentry and Halevi 2011b] next). Thus, the final step in Gentry’s blueprint is to squash the decryption circuit of the SWHE scheme, namely, transform the scheme into one with the same homomorphic capacity but with a decryption circuit that is simple enough to allow bootstrapping. Gentry [2009b] showed how to do this by adding a “hint”, namely, a large set of numbers with a secret sparse subset that sums to the original secret key, to the public key. Of course, the hint can be seen as useful information about the secret key, and the security of the scheme in the presence of the hint relies on a new “sparse subset sum” assumption (which, roughly speaking, can be thought of as saying that the hint is useless to a computationally-bounded adversary).

1.1. Efficiency of FHE

The efficiency of fully homomorphic encryption has been a (perhaps, the) big question following its invention. In this article, we are concerned with the *per-gate computation*

overhead of the FHE scheme, defined as the ratio between the time it takes to compute a circuit homomorphically on encrypted inputs to the time it takes to compute it on plaintext inputs.¹ Unfortunately, FHE schemes that follow Gentry's blueprint (some of which have actually been implemented [Coron et al. 2011; Gentry and Halevi 2011b]) have fairly poor performance: their per-gate computation overhead is $p(\lambda)$, a large polynomial in the security parameter. In fact, as we argue next, this penalty in performance seems somewhat inherent for schemes that follow this blueprint.

First, the complexity of (known approaches to) bootstrapping is inherently at least the complexity of decryption times the bit-length of the individual ciphertexts that are used to encrypt the bits of the secret key, the reason being that bootstrapping involves evaluating the decryption circuit *homomorphically*, that is, in the decryption circuit, each secret-key bit is replaced by a (large) ciphertext that encrypts that bit, and both the complexity of decryption and the ciphertext lengths must each be $\Omega(\lambda)$.

Second, the undesirable properties of known SWHE schemes conspire to ensure that the real cost of bootstrapping for FHE schemes that follow this blueprint is actually much worse than quadratic. Known FHE schemes start with an SWHE scheme that can evaluate polynomials of degree D (multiplicative depth $\log D$) securely only if the underlying lattice problem is hard to 2^D -approximate. To achieve hardness against 2^λ time adversaries, the lattice must have dimension $\Omega(D \cdot \lambda)$, because we have lattice algorithms in n dimensions that compute $2^{n/\lambda}$ -approximations of short vectors in time $2^{\tilde{O}(\lambda)}$. Moreover, the coefficients of the vectors used in the scheme have bit length $\Omega(D)$ to allow the ciphertext noise room to expand to 2^D . Therefore, the size of “fresh” ciphertexts (e.g., those that encrypt the bits of the secret key) is $\tilde{\Omega}(D^2 \cdot \lambda)$. Since the SWHE scheme must be bootstrappable, that is, capable of evaluating its own decryption function, D must exceed the degree of the decryption function. Typically, the degree of the decryption function is $\Omega(\lambda)$. Thus, overall, fresh ciphertexts have size $\tilde{\Omega}(\lambda^3)$. So, the real cost of bootstrapping—even if we optimistically assume that the “stale” ciphertext that needs to be refreshed can be decrypted in only $\Theta(\lambda)$ -time—is $\tilde{\Omega}(\lambda^4)$.

This analysis ignores a nice optimization by Stehlé and Steinfeld [2010] that uses Chernoff bounds to asymptotically reduce the decryption degree down to $O(\sqrt{\lambda})$. With this optimization, the per-gate computation of FHE schemes that follow the blueprint is $\tilde{\Omega}(\lambda^3)$.²

1.2. Recent Deviations from Gentry's Blueprint, and the Hope for Better Efficiency

Recently, Gentry and Halevi [2011a], and Brakerski and Vaikuntanathan [2011b], independently found very different ways to construct FHE without using the squashing step, and thus without the sparse subset sum assumption. These schemes are the first major deviations from Gentry's blueprint for FHE. Surprisingly, Brakerski and Vaikuntanathan [2011b] showed how to base security entirely on LWE (for sub-exponential approximation factors), avoiding reliance on ideal lattices.

From an efficiency perspective, however, these results are not a clear win over previous schemes. Both of the schemes still rely on the problematic aspects of Gentry's blueprint, namely, bootstrapping and an SWHE scheme with the undesirable

¹Other measures of efficiency, such ciphertext/key size and encryption/decryption time, are also important. In fact, the schemes we present in this article are very efficient in these aspects.

²We note that bootstrapping lazily, that is, applying the refresh procedure only at a $1/L$ fraction of the circuit levels for $L > 1$, cannot reduce the per-gate computation further by more than a logarithmic factor for schemes that follow this blueprint, since these SWHE schemes can evaluate only log-multiplicative depth before it becomes absolutely necessary to refresh, that is, $L = O(\log \lambda)$.

properties just discussed. Thus, their per-gate computation is still more than $\tilde{\Omega}(\lambda^4)$. Nevertheless, the techniques introduced in these recent constructions are very interesting and useful to us. In particular, we use the tools and techniques introduced by Brakerski and Vaikuntanathan [2011b] in an essential way to achieve remarkable efficiency gains.

An important, somewhat orthogonal question is the strength of assumptions underlying FHE schemes. All the schemes so far rely on the hardness of short vector problems on lattices with a subexponential approximation factor. Can we base FHE on the hardness of finding a polynomial approximation?

1.3. Our Results and Techniques

We leverage Brakerski and Vaikuntanathan's techniques [2011b] to achieve asymptotically very efficient FHE schemes. Also, we base security on lattice problems with *quasipolynomial* approximation factors. (All previous schemes relied on the hardness of problems with subexponential approximation factors.) In particular, we present the following two results.

- Assuming Ring LWE for an approximation factor exponential in L , we have a leveled FHE scheme that can evaluate L -level arithmetic circuits without using bootstrapping. The scheme has $\tilde{O}(\lambda \cdot L^3)$ per-gate computation (namely, quasilinear in the security parameter).
- Alternatively, assuming Ring LWE is hard for quasipolynomial factors, we have a leveled FHE scheme that uses bootstrapping as an optimization, where the per-gate computation (which includes the bootstrapping procedure) is $\tilde{O}(\lambda^2)$, independent of L .

We can alternatively base security on LWE, albeit with worse performance. We now sketch our main idea for boosting efficiency.

In the scheme of Brakerski and Vaikuntanathan [2011b], like ours, a ciphertext vector $\mathbf{c} \in R^n$ (where R is a ring, and n is the *dimension* of the vector) that encrypts a message m satisfies the decryption formula $m = [\lfloor \langle \mathbf{c}, \mathbf{s} \rangle \rfloor_q]_2$, where $\mathbf{s} \in R^n$ is the secret key vector, q is an odd modulus, and $[\cdot]_q : \mathbb{R} \rightarrow [-q/2, q/2)$ denotes the modular reduction function that reduces a real number x into the range $(-q/2, q/2)$. This is an abstract scheme that can be instantiated with either LWE or Ring LWE. In the LWE instantiation, R is the ring of integers mod q and n is a large dimension, whereas in the Ring LWE instantiation, R is the ring of polynomials over integers mod q and an irreducible $f(x)$, and the dimension $n = 2$.

We will call $\langle \mathbf{c}, \mathbf{s} \rangle$ the *noise* associated to ciphertext \mathbf{c} under key \mathbf{s} . Decryption succeeds as long as the magnitude of the noise stays smaller than $q/2$. Homomorphic addition and multiplication increase the noise in the ciphertext. The addition of two ciphertexts with noise at most B results in a ciphertext with noise at most $2B$, whereas multiplication results in a noise as large as B^2 .³ We will describe a *noise-management technique* that keeps the noise in check by reducing it after homomorphic operations, without bootstrapping.

The key technical tool we use for noise management is the *modulus-switching* technique developed by Brakerski and Vaikuntanathan [2011b]. Jumping ahead, we note that while they use modulus switching in one shot to obtain a small ciphertext (to which they then apply Gentry's bootstrapping procedure), we will use it (iteratively,

³The noise after multiplication is in fact a bit larger than B^2 due to the additional noise from the re-linearization process in Brakerski and Vaikuntanathan [2011a, 2011b]. For the purposes of this exposition, it is best to ignore this minor detail.

gradually) to keep the noise level essentially constant, while stingily sacrificing modulus size and gradually sacrificing the remaining homomorphic capacity of the scheme.

1.4. Modulus Switching

The essence of the modulus-switching technique is captured in the following lemma. In other words, the lemma says that an evaluator, who does not know the secret key \mathbf{s} but instead only knows a bound on its length, can transform a ciphertext \mathbf{c} modulo q into a different ciphertext modulo p , while preserving correctness, namely, $[\langle \mathbf{c}', \mathbf{s} \rangle]_p = [\langle \mathbf{c}, \mathbf{s} \rangle]_q \text{ mod } 2$. The transformation from \mathbf{c} to \mathbf{c}' involves simply scaling by (p/q) and rounding appropriately. Most interestingly, if \mathbf{s} is short and p is sufficiently smaller than q , the noise in the ciphertext actually decreases, namely, $|[\langle \mathbf{c}', \mathbf{s} \rangle]_p| < |[\langle \mathbf{c}, \mathbf{s} \rangle]_q|$.

This lemma is such an important centerpiece of our work that we will state and prove it right here.

LEMMA 1.1. *Let p and q be two odd moduli, and let \mathbf{c} be an integer vector. Define \mathbf{c}' to be the integer vector closest to $(p/q) \cdot \mathbf{c}$ such that $\mathbf{c}' = \mathbf{c} \bmod 2$. Then, for any \mathbf{s} with $|[\langle \mathbf{c}, \mathbf{s} \rangle]_q| < q/2 - (q/p) \cdot \ell_1(\mathbf{s})$, we have*

$$[\langle \mathbf{c}', \mathbf{s} \rangle]_p = [\langle \mathbf{c}, \mathbf{s} \rangle]_q \text{ mod } 2, \quad \text{and} \quad (1)$$

$$|[\langle \mathbf{c}', \mathbf{s} \rangle]_p| < (p/q) \cdot |[\langle \mathbf{c}, \mathbf{s} \rangle]_q| + \ell_1(\mathbf{s}), \quad (2)$$

where $\ell_1(\mathbf{s})$ is the ℓ_1 -norm of \mathbf{s} .

PROOF. For some integer k , we have $[\langle \mathbf{c}, \mathbf{s} \rangle]_q = \langle \mathbf{c}, \mathbf{s} \rangle - kq$. For the same k , let $e_p = \langle \mathbf{c}', \mathbf{s} \rangle - kp \in \mathbb{Z}$. Since $\mathbf{c}' = \mathbf{c}$ and $p = q \bmod 2$, we have $e_p = [\langle \mathbf{c}, \mathbf{s} \rangle]_q \text{ mod } 2$. Therefore, to prove the lemma, it suffices to prove that $e_p = [\langle \mathbf{c}', \mathbf{s} \rangle]_p$ and that it has small enough norm. We have $e_p = (p/q)[\langle \mathbf{c}, \mathbf{s} \rangle]_q + \langle \mathbf{c}' - (p/q)\mathbf{c}, \mathbf{s} \rangle$, and therefore $|e_p| \leq (p/q)|[\langle \mathbf{c}, \mathbf{s} \rangle]_q| + \ell_1(\mathbf{s}) < p/2$. The latter inequality implies $e_p = [\langle \mathbf{c}', \mathbf{s} \rangle]_p$. \square

Amazingly, this trick permits the evaluator to reduce the magnitude of the noise without knowing the secret key, and without bootstrapping. In other words, modulus switching gives us a very powerful and lightweight way to manage the noise in FHE schemes. In Brakerski and Vaikuntanathan [2011b], the modulus-switching technique is bundled into a “dimension reduction” procedure, and we believe it deserves a separate name and close scrutiny. It is also worth noting that our use of modulus switching does not require an “evaluation key”, in contrast to Brakerski and Vaikuntanathan [2011b].

1.5. Our New Noise Management Technique

At first, it may look like modulus switching is not a very effective noise management tool. If p is smaller than q , then of course modulus switching may reduce the magnitude of the noise, but it reduces the modulus size by essentially the same amount. In short, the ratio of the noise to the *noise ceiling* (the modulus size) does not decrease at all. Isn’t this ratio what dictates the remaining homomorphic capacity of the scheme, and how could potentially worsening (certainly not improving) this ratio do anything useful?

In fact, it’s not just the ratio of the noise to the noise ceiling that’s important. The absolute magnitude of the noise is also important, especially in multiplications. Suppose that $q \approx x^k$, and that you have two mod- q SWHE ciphertexts with noise of magnitude x . If you multiply them, the noise becomes x^2 . After four levels of multiplication, the noise is x^{16} . If you do another multiplication at this point, you reduce the ratio of the noise ceiling (i.e., q) to the noise level by a huge factor of x^{16} , that is, you reduce this

gap very fast. Thus, the actual magnitude of the noise impacts how fast this gap is reduced. After only $\log k$ levels of multiplication, the noise level reaches the ceiling.

Now, consider the following alternative approach. Choose a *ladder of gradually decreasing moduli* $\{q_i \approx q/x^i\}$ for $i < k$. After you multiply the two mod- q ciphertexts, switch the ciphertext to the smaller modulus $q_1 = q/x$. As the preceding lemma shows, the noise level of the new ciphertext (now with respect to the modulus q_1) goes from x^2 back down to x . (Let's suppose for now that $\ell_1(\mathbf{s})$ is small in comparison to x so that we can ignore it.) Now, when we multiply two ciphertexts (w.r.t. modulus q_1) that have noise level x , the noise again becomes x^2 , but then we switch to modulus q_2 to reduce the noise back to x . In short, each level of multiplication only reduces the ratio (noise ceiling)/(noise level) by a factor of x (not something like x^{16}). With this new approach, we can perform about k (not just $\log k$) levels of multiplication before we reach the noise ceiling. We have just increased (without bootstrapping) the number of multiplicative levels that we can evaluate by an exponential factor.

This exponential improvement is enough to achieve leveled FHE without bootstrapping. For any polynomial L , we can evaluate circuits of depth L . The performance of the scheme degrades with L (e.g., we need to set $q = q_0$ to have bit length proportional to L), but it degrades only polynomially with L .

Our main observation—the key to obtaining FHE without bootstrapping—is so simple that it is easy to miss and bears repeating: We get noise reduction automatically via modulus switching, and by carefully calibrating our ladder of moduli $\{q_i\}$ —one modulus for each circuit level—to be decreasing gradually, we can keep the noise level very small and essentially constant from one level to the next, while only gradually sacrificing the size of our modulus until the ladder is used up. With this approach, we can efficiently evaluate arbitrary polynomial-size arithmetic circuits without resorting to bootstrapping.

In terms of performance, this scheme trounces previous FHE schemes in terms of both asymptotic and concrete performance. Instantiated with ring-LWE, it can evaluate L -level arithmetic circuits with per-gate computation $\tilde{O}(\lambda \cdot L^3)$, that is, computation quasilinear in the security parameter. Since the ratio of the largest modulus (namely, $q \approx x^L$) to the noise (namely, x) is exponential in L , the scheme relies on the hardness of approximating short vectors to within an exponential in L factor.

1.6. Bootstrapping for Better Efficiency and Better Assumptions

In our FHE-without-bootstrapping scheme, the per-gate computation depends polynomially on the number of levels in the circuit that is being evaluated. While this approach is efficient (in the sense of polynomial time) for polynomial-size circuits, the per-gate computation may become undesirably high for very deep circuits, so we re-introduce bootstrapping as an optimization that makes the per-gate computation independent of the circuit depth, and that (if one is willing to assume circular security) allows homomorphic operations to be performed indefinitely without needing to specify in advance a bound on the number of circuit levels. The main idea is that to compute arbitrary polynomial-depth circuits, it is enough to compute the decryption circuit of the scheme homomorphically. Since the decryption circuit has depth $\approx \log \lambda$, the largest modulus we need has only $\text{polylog}(\lambda)$ bits, and therefore we can base security on the hardness of lattice problems with quasipolynomial factors. Since the decryption circuit has size $\tilde{O}(\lambda)$ for the RLWE-based instantiation, the per-gate computation becomes $\tilde{O}(\lambda^2)$ (independent of L). See Section 6 for details.

We then consider *batching* as an optimization. The idea behind batching is to pack multiple plaintexts into each ciphertext so that a function can be homomorphically

evaluated on multiple inputs with approximately the same efficiency as homomorphically evaluating it on one input.

An especially interesting case is *batching the decryption function* so that multiple ciphertexts (e.g., all of the ciphertexts associated to gates at some level in the circuit) can be bootstrapped simultaneously very efficiently. For circuits of large width (say, width λ), batched bootstrapping reduces the per-gate computation in the RLWE-based instantiation to $\tilde{O}(\lambda)$, independent of L . We give the details in Section 6.

1.7. Related and Subsequent Work

Prior to Gentry's construction, there were already a few interesting homomorphic encryption schemes that could be called "somewhat homomorphic", including that of Boneh et al. [2005] (that evaluates quadratic formulas using bilinear maps), Melchor et al. [2010] (that evaluates constant degree polynomials using lattices), and Ishai and Paskin [2007] (that evaluates branching programs).

Our work has inspired a number of follow-up works, building upon it and extending it in both theoretical and practical directions.

Our RLWE-based FHE scheme without bootstrapping requires only $\tilde{O}(\lambda \cdot L^3)$ per-gate computation, where L is the depth of the circuit being evaluated, while the bootstrapped version has only $\tilde{O}(\lambda^2)$ per-gate computation. For circuits of width $\Omega(\lambda)$, we can use batching to reduce the per-gate computation of the bootstrapped version by another factor of λ . In a follow-up work, Gentry et al. [2012b] showed how to reduce the per-gate evaluation overhead to polylogarithmic, making clever use of sorting networks to avoid packing/unpacking after every batched operation. They further suggested [Gentry et al. 2012a] choices of parameters so as to optimize the implementation of RLWE-based schemes. However, the polylogarithmic factors in these constructions are still too large to offer improvement for any "reasonable" value of the security parameter. One future direction toward a truly practical scheme is to tighten up these polylogarithmic factors considerably.

Using the ideas from this work in an essential way, Gentry et al. [2012c] implemented a variant of our scheme and presented benchmarks for homomorphic evaluation of the AES function.

Brakerski [2012] showed how to achieve comparable performance to ours without using modulus switching for arbitrary values of "initial" q . This is achieved by scaling the ciphertexts at the beginning of time rather than doing so after every operation. Gentry et al. [2012a] showed how to perform ideal reduction in the general case (and not just for the rings that we consider), significantly extending our ideas from Section 3.

More recently, Halevi and Shoup [2013] implemented our scheme together with the optimizations proposed by Gentry et al. in an open-source homomorphic encryption library called *HElib*.

2. PRELIMINARIES

2.1. Basic Notation

In our construction, we will use a ring R , whose elements we write in the lower case (e.g., $r \in R$). The reader should think of concrete instantiations of the ring to be either the ring of rational integers \mathbb{Z} or the polynomial ring $\mathbb{Z}[x]/(x^d + 1)$ for d a power of 2. When R is the ring of rational integers \mathbb{Z} , the notation $\|r\|$ refers to the absolute magnitude of r . When R is the polynomial ring $\mathbb{Z}[x]/(x^d + 1)$, $r \in R$ is a polynomial of degree at most $d - 1$, and the notation $\|r\|$ refers to the Euclidean norm of the coefficient vector of r . We also refer to $\|r\|$ as the *length* of r .

We write vectors in boldface (e.g., $\mathbf{v} \in R^n$). The notation $\mathbf{v}[i]$ refers to the i th coefficient of \mathbf{v} . We write the scalar product of $\mathbf{u}, \mathbf{v} \in R^n$ as $\langle \mathbf{u}, \mathbf{v} \rangle = \sum_{i=1}^n \mathbf{u}[i] \cdot \mathbf{v}[i] \in R$.

We define

$$\gamma_R := \max\{\|a \cdot b\| / \|a\| \|b\| : a, b \in R\}$$

to be the expansion factor of R . It is easy to see that for $R = \mathbb{Z}$, $\gamma_R = 1$ and for $R = \mathbb{Z}[x]/(x^d + 1)$, $\gamma_R \leq \sqrt{d}$ by the Cauchy-Schwarz inequality. We provide a proof for completeness.

PROPOSITION 2.1. *Let $R = \mathbb{Z}[x]/(x^d + 1)$. Then, $\gamma_R = \sqrt{d}$.*

PROOF. Let $a, b \in R$. Then, each coefficient of $a \cdot b$ is a scalar product of coefficients of a and b (or their rotations), modulo the sign. Thus, each coefficient has magnitude at most $\|a\| \cdot \|b\|$ by the Cauchy-Schwarz inequality. So, $\|a \cdot b\| \leq \sqrt{d} \cdot \|a\| \cdot \|b\|$. \square

Although working with the canonical embedding of $\mathbb{Z}[x]/(x^d + 1)$ provides a tighter way of handling the geometry of cyclotomic rings [Lyubashevsky et al. 2010] and eliminates the need to work with the expansion factor, we prefer to work with the coefficient representation of $r \in \mathbb{Z}[x]/(x^d + 1)$ in this exposition, as it is simple to describe and suffices for our asymptotic results.

For an integer q , we use R_q to denote R/qR . For $a \in R$, we use the notation $[a]_q$ to refer to $a \bmod q$, with coefficients reduced into the range $(-q/2, q/2]$.

Sometimes we will abuse notation and use R_2 to denote the set of R -elements with binary coefficients. For example, when $R = \mathbb{Z}$, R_2 may denote $\{0, 1\}$, and when R is a polynomial ring, R_2 may denote those polynomials whose coefficients are either 0 or 1. We use $R_{q,d}$ when we also want to specify the degree of the polynomial associated to R . When it is obvious that q is not a power of two, we will use $\lceil \log q \rceil$ to denote $1 + \lfloor \log q \rfloor$.

Finally, we define what it means for a distribution (ensemble) to be bounded.

Definition 2.2 (B-Bounded Distributions). A distribution ensemble $\{\chi_\lambda\}_{\lambda \in \mathbb{N}}$, supported over some normed space, is B -bounded for $B = B(\lambda)$ if

$$\Pr_{e \leftarrow \chi_\lambda} [\|e\| > B] = 0.$$

2.2. Homomorphic Encryption

Our definitions here largely follow the exposition of Brakerski and Vaikuntanathan [2011b]. We do take the liberty to slightly modify the algorithmic interface to better fit our constructions, and we urge even the knowledgeable reader to go over them so as to avoid confusion in interpreting our results.

We start with a generic definition of homomorphic encryption that captures both partially and fully homomorphic schemes. A homomorphic (public-key) encryption scheme is a tuple of polynomial-time algorithms ($\text{HE}.\text{Setup}$, $\text{HE}.\text{SecretKeyGen}$, $\text{HE}.\text{PublicKeyGen}$, $\text{HE}.\text{Enc}$, $\text{HE}.\text{Dec}$, $\text{HE}.\text{Eval}$) (although, in our exposition, we will start by describing a regular, non-homomorphic, public-key encryption scheme which has the same syntax as the preceding except for the homomorphic evaluation algorithm.)

In this work, the message space \mathcal{M} of the encryption schemes will always be some ring $R_{\mathcal{M}}$, and the functions to be evaluated will be represented as arithmetic circuits over this ring, composed of addition and multiplication gates.

The syntax of these algorithms is as follows.

— *Setup*. The algorithm $\text{params} \leftarrow \text{HE}.\text{Setup}(1^\lambda)$ takes as input the security parameter λ and outputs the global parameters of the encryption scheme. Jumping ahead, we remark that in our LWE-, and RLWE-based constructions, this will include a

number of parameters, such as a dimension n , a modulus q , an error distribution χ , and so on.

— *Key Generation.* Key generation consists of a pair of probabilistic polynomial-time algorithms. The secret key generation algorithm $sk \leftarrow \text{HE.SecretKeyGen}(\text{params})$ takes the global parameters and outputs a secret decryption key sk . The public key generation algorithm $pk \leftarrow \text{HE.PublicKeyGen}(\text{params}, sk)$ takes as input the global parameters and the secret key, and outputs the public key.

All our algorithms take the global parameters params as input, even when this is not explicitly specified.

— *Encryption.* The probabilistic encryption algorithm $c \leftarrow \text{HE.Enc}(pk, \mu)$ takes the public key pk and a message $\mu \in R_{\mathcal{M}}$ and outputs a ciphertext c .

— *Decryption.* The decryption algorithm $\mu^* \leftarrow \text{HE.Dec}(sk, c)$ takes the secret key sk and a ciphertext c and outputs a message $\mu^* \in R_{\mathcal{M}}$.

— *Homomorphic Evaluation.* The evaluation algorithm $c_f \leftarrow \text{HE.Eval}(pk, f, c_1, \dots, c_\ell)$ takes the public key pk , a function $f : R_{\mathcal{M}}^\ell \rightarrow R_{\mathcal{M}}$ (represented as an arithmetic circuit over $R_{\mathcal{M}}$), and a set of ℓ ciphertexts c_1, \dots, c_ℓ , and outputs a ciphertext c_f .

The only security notion we consider in this work is semantic security, namely, security with respect to passive adversaries [Goldwasser and Micali 1984]. We use its widely known formulation as IND-CPA security, defined as follows.

Definition 2.3 (CPA Security). A scheme HE is IND-CPA secure if for any polynomial time adversary $(\mathcal{A}, \mathcal{B})$, there is a negligible function negl whose advantage, defined as follows, is negligible:

$$\begin{aligned} \text{Adv}_{\text{cpa}}[\mathcal{A}, \mathcal{B}] := & |\Pr[(m_0, m_1) \leftarrow \mathcal{A}(pk); c_0 \leftarrow \text{HE.Enc}(pk, m_0) : \mathcal{B}(pk, c_0) = 1] \\ & - \Pr[(m_0, m_1) \leftarrow \mathcal{A}(pk); c_1 \leftarrow \text{HE.Enc}(pk, m_1) : \mathcal{B}(pk, c_1) = 1]| = \text{negl}(\lambda), \end{aligned}$$

where the probabilities are over the random choice of $sk \leftarrow \text{HE.PublicKeyGen}(\text{params})$ and $pk \leftarrow \text{HE.PublicKeyGen}(\text{params}, sk)$.

We move on to define the homomorphism property. Note that we do not define the *correctness* of the scheme as a separate property, but rather (some form of) correctness will follow from our homomorphism properties.

We start by defining \mathcal{F} -homomorphism, which is homomorphism with respect to a specified class \mathcal{F} of functions. This notion is sometimes also referred to as “somewhat homomorphism”. (The representation of the function f is an important issue. Since the representation can vary between schemes, we leave this issue outside of this syntactic definition. We remark, however, that in this work, f will be represented by an arithmetic circuit over $\text{GF}(2)$.)

Definition 2.4 (\mathcal{F} -Homomorphism). Let $\mathcal{F} = \{\mathcal{F}_\lambda\}_{\lambda \in \mathbb{N}}$ be a class of functions (together with their respective representations). A scheme HE is \mathcal{F} -homomorphic (or, homomorphic for the class \mathcal{F}) if for any sequence of functions $f_\lambda \in \mathcal{F}_\lambda$ and respective inputs $\mu_1, \dots, \mu_\ell \in \mathcal{M}$ (where \mathcal{M} is the message space, and $\ell = \ell(\lambda)$), it holds that

$$\Pr[\text{HE.Dec}(sk, \text{HE.Eval}(pk, f, c_1, \dots, c_\ell)) \neq f(\mu_1, \dots, \mu_\ell)] = \text{negl}(\lambda),$$

where the global parameters $\text{params} \leftarrow \text{HE.Setup}(1^\lambda)$, the secret and public keys are generated as $sk \leftarrow \text{HE.SecretKeyGen}(\text{params})$, $pk \leftarrow \text{HE.PublicKeyGen}(\text{params}, sk)$, and the ciphertexts $c_i \leftarrow \text{HE.Enc}(pk, \mu_i)$.

We point out two important properties that is preceding definition does not require. First of all, we do not require that the ciphertexts c_i are decryptable themselves, only that they become decryptable after homomorphic evaluation. While this may seem strange at first, this notion of somewhat homomorphism is all that is really required in order to bootstrap into full homomorphism and it also makes our schemes easier to describe. Note that one can always perform a “blank” homomorphic operation and then decrypt, so functionality is not hurt.

Second, we do not require that the output of $\text{HE}.\text{Eval}$ undergo additional homomorphic evaluation. This is termed “1-hop homomorphism” in Gentry et al. [2010].

Before we define full homomorphism, let us define the notion of *compactness*.

Definition 2.5 (Compactness). A homomorphic scheme HE is compact if there exists a polynomial $s = s(\lambda)$ such that the output length of $\text{HE}.\text{Eval}(\dots)$ is at most s -bits long (regardless of f or the number of inputs).

Note that an \mathcal{F} -homomorphic scheme is not necessarily compact.

We give the minimal definition of fully homomorphic encryption, which suffices for most applications.

Definition 2.6 (Fully Homomorphic Encryption). A scheme HE is fully homomorphic if it is both compact and homomorphic for the class of all arithmetic circuits over $GF(2)$.

As in the definition of \mathcal{F} homomorphism, one could require that the outputs of $\text{HE}.\text{Eval}$ could again be used as inputs for homomorphic evaluation (multihop homomorphism). Indeed, any bootstrappable scheme (see Section 2.3), including ours, has this additional property. However, due to the complexity of the formal definition in this case, we refrain from presenting a formal definition.

An important relaxation of fully homomorphic encryption is the following.

Definition 2.7 (Leveled Fully Homomorphic Encryption). A leveled, fully homomorphic encryption scheme is a homomorphic scheme where the $\text{HE}.\text{SecretKeyGen}$ gets an additional input 1^L (now $sk \leftarrow \text{HE}.\text{SecretKeyGen}(\text{params}, 1^L)$), and the resulting scheme is homomorphic for all depth- L binary arithmetic circuits. The bound $s(\lambda)$ on the ciphertext length must remain independent of L .

Most of this article will focus on the construction of a *leveled* fully homomorphic scheme in the sense that the parameters of the scheme depend (polynomially) on the depth of the circuits that the scheme is capable of evaluating. In most cases, the only parameter of the scheme that becomes dependent on L is the bit-length of the public key pk .

2.3. Gentry’s Bootstrapping Technique

In this section, we formally define the notion of a bootstrappable encryption scheme and present Gentry’s bootstrapping theorem [Gentry 2009a, 2009b] which implies that a bootstrappable scheme can be converted into a fully homomorphic one.

Definition 2.8 (Bootstrappable Encryption Scheme). Let HE be \mathcal{F} -homomorphic, and let f_{add} and f_{mult} be the augmented decryption functions of the scheme defined as

$$f_{\text{add}}^{c_1, c_2}(s) = \text{HE}.\text{Dec}_s(c_1) \text{ XOR } \text{HE}.\text{Dec}_s(c_2) \text{ and } f_{\text{mult}}^{c_1, c_2}(s) = \text{HE}.\text{Dec}_s(c_1) \text{ AND } \text{HE}.\text{Dec}_s(c_2),$$

where c_1, c_2 are either properly encrypted ciphertexts of the scheme or outputs of the homomorphic evaluation function, applied to such.

Then HE is bootstrappable if

$$\{f_{\text{add}}^{c_1, c_2}, f_{\text{mult}}^{c_1, c_2}\}_{c_1, c_2} \subseteq \mathcal{F}.$$

Namely, the scheme can homomorphically evaluate f_{add} and f_{mult} .

We describe two variants of Gentry's bootstrapping theorem. The first implies leveled, fully homomorphic encryption but requires no additional assumption, where the second makes an additional (*weak*) *circular security* assumption and achieves the stronger (non-leveled) variant of Definition 2.6.

The first variant follows.

THEOREM 2.9 [GENTRY 2009A, 2009B]. *If there exists a bootstrappable encryption scheme, then there exists a leveled, fully homomorphic encryption scheme as per Definition 2.7.*

Specifically, the leveled homomorphic scheme is such that only the length of the evaluation key depends on the level L . All other parameters of the scheme are distributed identically regardless of the value of L .

For the second variant, we need to define circular security.

Definition 2.10 (Weak Circular Security). A public key encryption scheme $(\text{KeyGen}, \text{Enc}, \text{Dec})$ is weakly circular secure if it is IND-CPA secure, even for an adversary with auxiliary information containing encryptions of all secret key bits: $\{\text{Enc}_{pk}(sk[i])\}_i$.

Namely, no polynomial-time adversary can distinguish an encryption of 0 from an encryption of 1, even given the additional information.

We can now state the second theorem.

THEOREM 2.11 [GENTRY 2009A, 2009B]. *If there exists a bootstrappable scheme that is also weakly circular secure, then there exists a fully homomorphic encryption scheme as per Definition 2.6.*

Finally, we want to make a statement regarding the ciphertext length of a bootstrapped scheme. The transformations in Gentry [2009a, 2009b] guarantees something stronger than what we claimed in Lemmas 2.9 and 2.11. Namely, starting from a bootstrappable scheme HE, one can construct a (leveled) FHE scheme FHE whose encryption algorithm FHE.Enc and evaluation algorithm FHE.Eval produce ciphertexts of the same length as HE.Eval (regardless of the length of the ciphertext produced by HE.Enc).

2.4. Learning with Errors Problem

The LWE problem was introduced by Regev [2005] as a generalization of “learning parity with noise”. For positive integers n and $q \geq 2$, a vector $\mathbf{s} \in \mathbb{Z}_q^n$, and a probability distribution χ on \mathbb{Z}_q , let $A_{\mathbf{s}, \chi}$ be the distribution obtained by choosing a vector $\mathbf{a} \leftarrow \mathbb{Z}_q^n$ uniformly at random and a noise term $e \leftarrow \chi$, and outputting $(\mathbf{a}, \langle \mathbf{a}, \mathbf{s} \rangle + e) \in \mathbb{Z}_q^n \times \mathbb{Z}_q$. A formal definition follows.

Definition 2.12 (LWE). For an integer $q = q(n)$ and an error distribution $\chi = \chi(n)$ over \mathbb{Z}_q , the learning with errors problem $\text{LWE}_{n, m, q, \chi}$ is defined as follows: Given m independent samples from $A_{\mathbf{s}, \chi}$ (for some $\mathbf{s} \in \mathbb{Z}_q^n$), output \mathbf{s} with noticeable probability.

The (average-case) decision variant of the LWE problem, denoted $\text{DLWE}_{n, m, q, \chi}$, is to distinguish (with nonnegligible advantage) m samples chosen according to $A_{\mathbf{s}, \chi}$ (for uniformly random $\mathbf{s} \leftarrow \mathbb{Z}_q^n$), from m samples chosen according to the uniform distribution over $\mathbb{Z}_q^n \times \mathbb{Z}_q$. We denote by $\text{DLWE}_{n, q, \chi}$ the variant where the adversary gets oracle access to $A_{\mathbf{s}, \chi}$ and is not a-priori bounded in the number of samples.

For an algorithm \mathcal{B} and security parameter λ , we denote

$$\text{DLWE}_{n,q,\chi} \text{Adv}[\mathcal{B}] := |\Pr[\mathcal{B}^{A_{\mathbf{s},\chi}}(1^\lambda) = 1] - \Pr[\mathcal{B}^{U(\mathbb{Z}_q^n \times \mathbb{Z}_q)}(1^\lambda) = 1]|.$$

For cryptographic applications, we are primarily interested in the average case decision problem DLWE, where $\mathbf{s} \leftarrow \mathbb{Z}_q^n$. There are known quantum [Regev 2005] and classical [Peikert 2009] reductions between $\text{DLWE}_{n,m,q,\chi}$ and approximating short vector problems in worst-case lattices. Specifically, these reductions take χ to be (discretized versions of) the Gaussian distribution. These distributions can easily be made B -bounded for an appropriate B by rejection sampling without effecting the validity of the reduction. Since the exact distribution χ does not matter for our results, we state a corollary of the results of Regev [2005] and Peikert [2009] in terms of the bound on the distribution.

COROLLARY 2.13 [REGEV 2005; PEIKERT 2009]. *Let $q = q(n) \in \mathbb{N}$ be a product of co-prime numbers $q = \prod q_i$ such that for all i , $q_i = \text{poly}(n)$, and let $B \geq n$. Then there exists an efficiently sampleable B -bounded distribution χ such that if there is an efficient algorithm that solves the (average-case) $\text{DLWE}_{n,q,\chi}$ problem. Then there exists an efficient algorithm that approximates short vector problems in n -dimensional lattices to within a $\tilde{O}(n\sqrt{n} \cdot q/B)$ factor. Specifically, the following.*

- *There is a quantum algorithm that solves $\text{SIVP}_{\tilde{O}(n\sqrt{n} \cdot q/B)}$ and $\text{gapSVP}_{\tilde{O}(n\sqrt{n} \cdot q/B)}$ on any n -dimensional lattice, and runs in time $\text{poly}(n)$.*
- *There is a classical algorithm that solves the ζ -to- γ decisional shortest vector problem $\text{gapSVP}_{\zeta,\gamma}$, where $\gamma = \tilde{O}(n\sqrt{n} \cdot q/B)$, and $\zeta = \tilde{O}(q\sqrt{n})$, on any n -dimensional lattice, and runs in time $\text{poly}(n)$.*

We refer the reader to Regev [2005] and Peikert [2009] for the formal definition of these lattice problems, as they have no direct connection to this work. We only note here that the best-known algorithms for these problems run in time nearly exponential in the dimension n [Ajtai et al. 2001; Micciancio and Voulgaris 2010]. More generally, the best algorithms that approximate these problems to within a factor of 2^k run in time $2^{\tilde{O}(n/k)}$ [Schnorr 1987].

Applebaum et al. [2009] showed that if LWE is hard for the preceding distribution of \mathbf{s} , then it is also hard when \mathbf{s} 's coefficients are sampled according to the noise distribution χ .

2.5. Ring Learning with Errors Problem

The ring learning with errors (RLWE) problem was introduced by Lyubashevsky et al. [2010]. We will use a simplified special-case version of the problem that is easier to work with [Regev 2010; Brakerski and Vaikuntanathan 2011a].

Definition 2.14 (RLWE). For security parameter λ , let $f(x) = x^d + 1$, where $d = d(\lambda)$ is a power of 2. Let $q = q(\lambda) \geq 2$ be an integer. Let $R = \mathbb{Z}[x]/(f(x))$ and let $R_q = R/qR$. Let $\chi = \chi(\lambda)$ be a distribution over R . The $\text{RLWE}_{d,q,\chi}$ problem is to distinguish the following two distributions: In the first distribution, one samples (a_i, b_i) uniformly from R_q^2 . In the second distribution, one first draws $s \leftarrow R_q$ uniformly and then samples $(a_i, b_i) \in R_q^2$ by sampling $a_i \leftarrow \mathbb{R}_q$ uniformly, $e_i \leftarrow \chi$, and setting $b_i = a_i \cdot s + e_i$. The $\text{RLWE}_{d,q,\chi}$ assumption is that the $\text{RLWE}_{d,q,\chi}$ problem is infeasible.

The RLWE problem is useful, because the well-established shortest vector problem (SVP) over ideal lattices can be reduced to it, specifically as follows.

THEOREM 2.15 (LYUBASHEVSKY ET AL. [2010]). *For any d that is a power of 2, ring $R = \mathbb{Z}[x]/(x^d + 1)$, prime integer $q = q(d) = 1 \bmod 2d$, and $B = \omega(d\sqrt{\log d})$, there is an efficiently samplable distribution χ that outputs elements of R of length at most B with overwhelming probability, such that if there exists an efficient algorithm that solves $\text{RLWE}_{d,q,\chi}$, then there is an efficient quantum algorithm for solving $f(d) \cdot (q/B)$ -approximate worst-case SVP for ideal lattices over R , for every super-polynomial factor $f(d) = d^{\omega(1)}$.*

Typically, to use RLWE with a cryptosystem, one chooses the noise distribution χ according to a Gaussian distribution, where vectors sampled according to this distribution have length only $\text{poly}(d)$ with overwhelming probability. This Gaussian distribution may need to be ellipsoidal for certain reductions to go through [Lyubashevsky et al. 2010]. It has been shown for RLWE that one can equivalently assume that s is alternatively sampled from the noise distribution χ [Lyubashevsky et al. 2010].

2.6. General Learning with Errors Problem

The learning with errors (LWE) problem and the ring learning with errors (RLWE) problem are syntactically identical, aside from using different rings (the ring of rational integers versus a polynomial ring) and different vector dimensions over those rings ($n = \text{poly}(\lambda)$ for LWE, but $n = 1$ for RLWE). To simplify our presentation, we define a *General Learning with Errors (GLWE) Problem* and describe a single GLWE-based FHE scheme, rather than presenting essentially the same scheme twice, once for each of our two concrete instantiations.

Definition 2.16 (GLWE). For security parameter λ , let $n = n(\lambda)$ be an integer dimension, let $f(x) = x^d + 1$ where $d = d(\lambda)$ is a power of 2, let $q = q(\lambda) \geq 2$ be a prime integer, let $R = \mathbb{Z}[x]/(f(x))$ and $R_q = R/qR$, and let $\chi = \chi(\lambda)$ be a distribution over R . The $\text{GLWE}_{n,f,q,\chi}$ problem is to distinguish the following two distributions: In the first distribution, one samples (\mathbf{a}_i, b_i) uniformly from R_q^{n+1} . In the second distribution, one first draws $\mathbf{s} \leftarrow R_q^n$ uniformly and then samples $(\mathbf{a}_i, b_i) \in R_q^{n+1}$ by sampling $\mathbf{a}_i \leftarrow R_q^n$ uniformly, $e_i \leftarrow \chi$, and setting $b_i = \langle \mathbf{a}_i, \mathbf{s} \rangle + e_i$. The $\text{GLWE}_{n,f,q,\chi}$ assumption is that the $\text{GLWE}_{n,f,q,\chi}$ problem is infeasible.

LWE is simply GLWE instantiated with $d = 1$. RLWE is GLWE instantiated with $n = 1$. Interestingly, as far as we know, instances of GLWE between these extremes have not been explored. One would suspect that GLWE is hard for any (n, d) such that $n \cdot d = \Omega(\lambda \log(q/B))$, where B is a bound (with overwhelming probability) on the length of elements output by χ . For fixed $n \cdot d$, perhaps GLWE gradually becomes harder as n increases (if it is true that general lattice problems are harder than ideal lattice problems), whereas increasing d is probably often preferable for efficiency.

If q is much larger than B , the associated GLWE problem is believed to be easier (i.e., there is less security). Previous FHE schemes required q/B to be subexponential in n or d to give room for the noise to grow as homomorphic operations (especially multiplication) are performed. In our FHE scheme without bootstrapping, q/B will be exponential in the number of circuit levels to be evaluated. However, since the decryption circuit can be evaluated in logarithmic depth, the bootstrapped version of our scheme will only need q/B to be quasipolynomial, and we thus base security on lattice problems for quasipolynomial approximation factors.

By the GLWE assumption, the distribution $\{(\mathbf{a}_i, \langle \mathbf{a}_i, \mathbf{s} \rangle + t \cdot e_i)\}$ is computationally indistinguishable from uniform for any t relatively prime to q . This fact will be

convenient for encryption, where, for example, a message m may be encrypted as $(\mathbf{a}, \langle \mathbf{a}, \mathbf{s} \rangle + 2e + m)$, and this fact can be used to argue that the second component of this message is indistinguishable from random (as shown in, e.g., [Brakerski and Vaikuntanathan 2011b].)

2.7. A Useful Lemma

We will use the following lemma that converts Boolean circuits into arithmetic circuits of roughly the same size and depth, in Section 6.

LEMMA 2.17. *Let $f : \{0, 1\}^n \rightarrow \{0, 1\}$ be a function. And let C be an arithmetic circuit over \mathbb{Z}_2 with $\{+, \times\}$ gates of fan-in 2 that computes f . Denote the size of C by s and its depth by d . Then there exists an arithmetic circuit C' with $\{+, -, \times\}$ gates of fan-in 2, of size $\leq 3s$ and depth $\leq 2d$, that computes f over any ring \mathbb{Z}_p , where $p \geq 2$ (but not necessarily prime). Further, C' can be efficiently computed from C .*

PROOF. The circuit C' will be defined by applying the following transformation on C . Every \times gate in C is left unchanged in C' . Any $+$ gate with input wires a, b in C is replaced by the gadget $a + b - ab$. Note that this gadget has size 3 and depth 2; therefore, the size of C' is at most $3s$ and its depth is at most $2d$.

The correctness of the construction follows, since for all $a, b \in \{0, 1\}$ and for any integer $p \geq 2$, the following holds.

- (1) $[a \cdot b]_2 = [a \cdot b]_p$.
- (2) $[a + b]_2 = [a + b - a \cdot b]_p$.

A straightforward inductive argument shows that for all $x_1, \dots, x_n \in \{0, 1\}$, it holds that $[C(x_1, \dots, x_n)]_2 = [C'(x_1, \dots, x_n)]_p$ for all p , and the result follows. \square

3. TRADING OFF DEGREE FOR DIMENSION IN GLWE

In this section, we present a new technique, termed *ring switching* in follow-up work [Gentry et al. 2012a], which allows trading off the parameters d and n of the GLWE problem. In particular, we show that these parameters can be traded off subject to the invariant that $n \cdot d$ remains constant. We note that, as per Definition 2.16, d must be a power of 2. As previously mentioned, this technique was generalized and extended in Gentry et al. [2012a].

Before moving on to the technical details, let us provide some motivation for trading off the degree and dimension. We show in the next section how to devise homomorphic encryption scheme with increased homomorphic capacity by increasing the degree of the underlying ring. However, this will incur an increase in the decryption complexity of the scheme, rendering it noncompact (see Definition 2.5). This is solved by first trading off the degree for dimension, as described, next, and then using the dimension reduction technique of Brakerski and Vaikuntanathan [2011b] to reduce the dimension, which will bring the final ciphertext to the same GLWE parameters regardless of the extent of the homomorphic evaluation.

We proceed by describing our technique. Namely, we show that there is an interplay between the dimension n of a GLWE problem and the degree d of the modulus polynomial. We show that an $GLWE_{n, x^d+1, q, \chi}$ ciphertext can be efficiently broken into two $GLWE_{2n, x^{d/2}+1, q, \chi}$ ciphertexts. We slightly deviate from the notation in the body of the article and denote $GLWE_{n, d, q, \chi}$ to denote $GLWE_{n, x^d+1, q, \chi}$ (recall that d is always a power of 2). We further denote $R_{q,d} = \mathbb{Z}[x]/(x^d + 1)$.

Decomposing Ring Elements. We begin by presenting a formal decomposition of elements from $R_{q,d}$ into elements of $R_{q,d/2}$. We show that each element $a = a(x) \in R_{q,d}$

can be represented using $a^{(\text{even})}, a^{(\text{odd})} \in R_{q,d/2}$. Recalling that an element in $R_{q,d}$ is a polynomial with d coefficients over \mathbb{Z}_q , the task seems very simple. We embed half of the coefficients of the polynomial a as coefficients of $a^{(\text{even})}$ and the other half as coefficients of $a^{(\text{odd})}$. However, in order to preserve an algebraic structure over the new elements, it is critical that the coefficients are divided between them in a special way.

Specifically, we define $a^{(\text{even})}, a^{(\text{odd})}$ to be the elements of $R_{q,d/2}$ for which

$$a(x) = a^{(\text{even})}(x^2) + x \cdot a^{(\text{odd})}(x^2).$$

In other words, $a^{(\text{even})}$ assumes the even coefficients of a , and $a^{(\text{odd})}$ the odd ones.

To see that the algebraic properties are preserved, suppose we have an equation $c(x) = a(x) \cdot b(x)$ over $R_{q,d}$, then it holds that

$$\begin{aligned} & c^{(\text{even})}(x^2) + xc^{(\text{odd})}(x^2) \\ &= (a^{(\text{even})}(x^2) + xa^{(\text{odd})}(x^2)) \cdot (b^{(\text{even})}(x^2) + xb^{(\text{odd})}(x^2)) \\ &= a^{(\text{even})}(x^2) \cdot b^{(\text{even})}(x^2) + x^2 a^{(\text{odd})}(x^2) \cdot b^{(\text{odd})}(x^2) \\ &\quad + x[a^{(\text{even})}(x^2) \cdot b^{(\text{odd})}(x^2) + a^{(\text{odd})}(x^2) \cdot b^{(\text{even})}(x^2)]. \end{aligned}$$

Noting that the parity of a power of x cannot change by reducing modulo $x^d + 1$ (since d is a power of 2 and thus always even), it follows that we can separate odd and even powers in the previous expression:

$$\begin{aligned} c^{(\text{even})}(x^2) &= a^{(\text{even})}(x^2) \cdot b^{(\text{even})}(x^2) + x^2 a^{(\text{odd})}(x^2) \cdot b^{(\text{odd})}(x^2), \\ c^{(\text{odd})}(x^2) &= a^{(\text{even})}(x^2) \cdot b^{(\text{odd})}(x^2) + a^{(\text{odd})}(x^2) \cdot b^{(\text{even})}(x^2). \end{aligned}$$

These equations are still over the ring $R_{q,d}$. In order to switch down to the ring $R_{q,d/2}$, we consider the following fact. In general, if $w(x^2) = u(x^2) \cdot v(x^2)$ over $R_{q,d}$, that is, modulo q and $x^d + 1$, then it holds that $w(x) = u(x) \cdot v(x)$ over $R_{q,d/2}$, that is, modulo q and $x^{d/2} + 1$. This follows syntactically by replacing x^2 everywhere with x .

Therefore, we have that the following holds over $R_{q,d/2}$:

$$\begin{aligned} c^{(\text{even})}(x) &= a^{(\text{even})}(x) \cdot b^{(\text{even})}(x) + xa^{(\text{odd})}(x) \cdot b^{(\text{odd})}(x), \\ c^{(\text{odd})}(x) &= a^{(\text{even})}(x) \cdot b^{(\text{odd})}(x) + a^{(\text{odd})}(x) \cdot b^{(\text{even})}(x). \end{aligned}$$

Decomposing Vectors. . The preceding technique can be applied in a straightforward way to decompose *vectors* of ring elements. This is done by decomposing each of the elements in the vector, yielding an odd vector that corresponds to all of the odd components of the elements, and an even vector corresponding to the even components. In order to complete our claim, we show that the inner product of the original vector can be expressed using the inner products of the decomposed vectors.

Let $\mathbf{a}, \mathbf{s} \in R_{q,d}^n$, and let $\tilde{\mathbf{b}} = \langle \mathbf{a}, \mathbf{s} \rangle$. Applying the decomposition to each component as just described yields vectors $\mathbf{a}^{(\text{even})}, \mathbf{a}^{(\text{odd})}, \mathbf{s}^{(\text{even})}, \mathbf{s}^{(\text{odd})}, \mathbf{b}^{(\text{even})}, \mathbf{b}^{(\text{odd})} \in R_{q,d/2}^n$ such that

$$\begin{aligned} \tilde{\mathbf{b}}^{(\text{even})} &= \langle \mathbf{a}^{(\text{even})}, \mathbf{s}^{(\text{even})} \rangle + x \langle \mathbf{a}^{(\text{odd})}, \mathbf{s}^{(\text{odd})} \rangle = \langle (\mathbf{a}^{(\text{even})}, x\mathbf{a}^{(\text{odd})}), (\mathbf{s}^{(\text{even})}, \mathbf{s}^{(\text{odd})}) \rangle, \\ \tilde{\mathbf{b}}^{(\text{odd})} &= \langle \mathbf{a}^{(\text{even})}, \mathbf{s}^{(\text{odd})} \rangle + \langle \mathbf{a}^{(\text{odd})}, \mathbf{s}^{(\text{even})} \rangle = \langle (\mathbf{a}^{(\text{odd})}, \mathbf{a}^{(\text{even})}), (\mathbf{s}^{(\text{even})}, \mathbf{s}^{(\text{odd})}) \rangle. \end{aligned}$$

Indeed, the original inner product breaks into two inner products of a reduced ring but of vectors of double the dimension. To apply this to GLWE, we need to consider the noisy version, $\mathbf{b} = \langle \mathbf{a}, \mathbf{s} \rangle + \mathbf{e}$, which is handled exactly as before but with the \mathbf{e} also decomposing into $\mathbf{e}^{(\text{odd})}$ and $\mathbf{e}^{(\text{even})}$.

The preceding step can be applied recursively k times to trade off GLWE with parameters (n, d) , for GLWE with parameters $(n \cdot 2^k, d/2^k)$.

4. LEVELED FHE WITHOUT BOOTSTRAPPING FROM LWE

The plan of this section is to present our *leveled FHE without bootstrapping* construction in modular steps. First, we describe a plain LWE-based encryption scheme with no homomorphic operations. Next, we augment the plain scheme with variants of the relinearization and dimension reduction techniques of Brakerski and Vaikuntanathan [2011b]. Finally, in Section 4.4, we lay out our full-fledged construction of FHE without bootstrapping.

4.1. The Basic Encryption Scheme

We begin by presenting a basic encryption scheme with no homomorphic operations.

Let λ be the security parameter, representing our goal of achieving 2^λ -security against known attacks. ($\lambda = 100$ is, e.g., a reasonable value.) Let $R = R(\lambda)$ be a ring. Specifically, we will be interested in two instantiations:

- (1) $R = \mathbb{Z}$, which will give us a scheme based on (standard) LWE, and
- (2) $R = \mathbb{Z}[x]/f(x)$, where, for example, $f(x) = x^d + 1$ and $d = d(\lambda)$ is a power of 2, which will give us a scheme based on RLWE with the appropriate parameters.

Let the *dimension* $n = n(\lambda)$, the degree $d = d(\lambda)$, an odd positive integer modulus $q = q(\lambda)$, and a *noise distribution* $\chi = \chi(\lambda)$ over R be parameters of the system which come from the LWE (resp. RLWE) assumption. Let $R_q := R/qR$. For simplicity, assume for now that the plaintext space is $R_2 := R/2R$, though we will discuss the issue of supporting larger plaintext spaces in the sequel. In addition to the usual LWE (resp. RLWE) parameters, we will use an additional parameter $N = N(\lambda) = n \cdot \text{polylog}(q)$ which we will discuss following the description of the scheme.

Even though this only becomes important when we introduce homomorphic operations, we go ahead and stipulate here that the noise distribution χ will be supported over a set of ring elements with small norm, namely, the following.

- (1) In the LWE case, the support of χ will be $[-B, \dots, B]$ for some $B = B(\lambda)$. We will set B to be as small as possible (while maintaining security), and
- (2) In the RLWE case, the support of χ will be the set of all ring elements with norm at most $B = B(\lambda)$.

The Basic Encryption Scheme. We now present the plain encryption scheme, which is inspired by the schemes of Regev [2005] in the LWE case, or Lyubashevsky et al. [2010] in the RLWE case. In the following description, we describe both the LWE and RLWE schemes using a common framework, highlighting the salient differences whenever appropriate.

The scheme $E = (E.\text{Setup}, E.\text{SecretKeyGen}, E.\text{PublicKeyGen}, E.\text{Enc}, E.\text{Dec})$ works as follows.

- $E.\text{Setup}(1^\lambda)$. Determine the system parameters as a function of the security parameter λ . That is, set $d = d(\lambda)$ to be the degree of the ring, and $n = n(\lambda)$ to be the dimension. In particular, in the LWE-based instantiation, the underlying ring $R = \mathbb{Z}$ (with degree $d = 1$), n will be determined based on the security parameter, and we will work over \mathbb{Z}_q^n . In the RLWE-based instantiation, the underlying ring $R = \mathbb{Z}[x]/f(x)$, where $f(x)$ has degree $d = d(\lambda)$, the dimension $n = 1$, and we will work over $\mathbb{Z}_q[x]/f(x)$.

Set the odd modulus $q = q(\lambda)$, the noise distribution $\chi = \chi(\lambda)$, and $N = N(\lambda) = n \cdot \text{polylog}(q)$. Output $\text{params} = (R, d, n, q, \chi, N)$.

- $E.\text{SecretKeyGen}(\text{params})$. Sample $\mathbf{t} \leftarrow \chi^n$. Let

$$\mathbf{s} := (1, \mathbf{t}[1], \dots, \mathbf{t}[n]) \in R_q^{n+1}.$$

Output the secret key $sk = \mathbf{s}$.

- E.PublicKeyGen($params, sk$). Take as input the secret key $sk = \mathbf{s} = (\mathbf{1}, \mathbf{t})$ with $\mathbf{s}[0] = 1$ and $\mathbf{t} \in R_q^n$ and the system parameters $params$. Generate a matrix $\mathbf{B} \leftarrow R_q^{N \times n}$ uniformly at random and a column vector with “small” coefficients $\mathbf{e} \leftarrow \chi^N$. Set $\mathbf{b} := \mathbf{B}\mathbf{t} + 2\mathbf{e}$ and the public key as

$$\mathbf{A} := (\mathbf{b} \parallel -\mathbf{B}) \in R_q^{N \times (n+1)}$$

Remark. Observe that $\mathbf{A} \cdot \mathbf{s} = \mathbf{b} \cdot \mathbf{1} - \mathbf{B} \cdot \mathbf{t} = 2\mathbf{e}$, by the way we defined \mathbf{b} .

- E.Enc($params, pk, m$). To encrypt a message $m \in R_2$, set the row vector $\mathbf{m} := (m, 0, \dots, 0) \in R_q^{n+1}$, sample a column vector with small coefficients $\mathbf{r} \leftarrow R_2^N$, and output the ciphertext

$$\mathbf{c} := \mathbf{m} + \mathbf{r}^T \cdot \mathbf{A} \in R_q^{n+1}.$$

- E.Dec($params, sk, \mathbf{c}$). Take as input the secret key $sk = \mathbf{s} \in R_q^{n+1}$ and a ciphertext $\mathbf{c} \in R_q^{n+1}$. Output the message $m := [\lfloor \langle \mathbf{c}, \mathbf{s} \rangle \rfloor_q]_2$.

We now informally describe why the scheme is correct and secure. Decryption works correctly because

$$[\lfloor \langle \mathbf{c}, \mathbf{s} \rangle \rfloor_q]_2 = [\lfloor (\mathbf{m} + \mathbf{r}^T \mathbf{A}) \cdot \mathbf{s} \rfloor_q]_2 = [\lfloor [m + 2\mathbf{r}^T \mathbf{e}]_q \rfloor_2 = [m + 2\mathbf{r}^T \mathbf{e}]_2 = m,$$

where the third equality holds because we will ensure that \mathbf{e} and \mathbf{r} have small enough entries so that that the value $m + 2\mathbf{r}^T \mathbf{e}$ does not wrap around modulo q .

It is straightforward to base security on the LWE or RLWE assumptions, for appropriate parameters (see, e.g., [Brakerski and Vaikuntanathan 2011b; Lyubashevsky et al. 2010; Regev 2005]). We will now sketch the main ideas. First, note that if an attacker can distinguish the public key \mathbf{A} from a uniformly random matrix over $R_q^{N \times (n+1)}$, then the attacker can be used to solve the LWE (resp. RLWE) problem. Therefore, assuming the LWE (resp. RLWE) problem is hard, an attacker cannot efficiently distinguish the public key from a uniformly random matrix. Second, if \mathbf{A} was indeed chosen uniformly from $R_q^{N \times (n+1)}$, the encryption procedure generates ciphertexts that are statistically independent from m , either by the leftover hash lemma (in the case of LWE, or by Micciancio’s regularity lemma [Micciancio 2007] in the case of RLWE). Put together, the attacker has negligible advantage in guessing the message m .

For the LWE case, it suffices to take $N > 2n \log q$ [Regev 2005]. For RLWE, it does not necessarily work just to take $N > 2n \log q = 2 \log q$ due to subtle distributional issues. In particular, the problem is that R_q may have many zero divisors. Micciancio’s regularity lemma assures us that if $\mathbf{A} \in R_q^{N \times (n+1)}$ and $\mathbf{r} \in R_2^N$ are uniform, then $\mathbf{r}^T \mathbf{A}$ has negligible statistical distance from uniform when $N = \log(q \cdot \lambda^{\omega(1)})$.

To achieve 2^λ security against known lattice attacks, one must have $n \cdot d = \Omega(\lambda \cdot \log(q/B))$, where B is a bound on the length of the noise. Since n or d depends logarithmically on q , and since the distribution χ (and hence B) depends sublinearly on n or d , the distribution χ (and hence B) depends sublogarithmically on q . This dependence is weak, and one should think of the noise distribution as being essentially independent of q .

Generalizing to Larger Plaintext Spaces. The scheme just described transparently handles plaintext spaces larger than R_2 . For example, to encrypt a message \mathbf{m} from the plaintext domain R_t , where t is relatively prime to q , we simply change the encryption to $(\mathbf{a}, \langle \mathbf{a}, \mathbf{s} \rangle + t\mathbf{e} + m)$. Of course, this also requires a slightly different analysis of how the error in the ciphertext grows with homomorphic operations.

An Alternate Encryption Scheme Based on RLWE. While we think our description of encryption is useful in that it highlights the high-level similarity of LWE and RLWE, the distributional issues previously discussed make it more desirable, in practice, to use a slightly different approach to encryption in the RLWE setting. In particular, Lyubashevsky et al. [2010] streamline public key generation and encryption in the RLWE setting as follows.

- E.PublicKeyGen($params, sk$). As before, except $N = 1$.
- E.Enc($params, pk, m$). To encrypt a message $m \in R_2$, set $\mathbf{m} := (m, 0) \in R_q^2$, sample $r \leftarrow \chi$ and a row vector $\mathbf{e} \leftarrow \chi^2$. Output the ciphertext

$$\mathbf{c} := \mathbf{m} + 2 \cdot \mathbf{e} + r \cdot \mathbf{A} \in R_q^2.$$

The security of LPR encryption relies on RLWE: assuming RLWE, the public key \mathbf{A} is computationally indistinguishable from uniform in R_q^2 . Then, invoking the RLWE assumption once again, the two ring elements $m + a_1 \cdot r + e_1$ and $a_2 \cdot r + e_2$ of the ciphertext generated during encryption are pseudorandom.

In what follows, some of our schemes will invoke the function E.PublicKeyGen($params, sk, N$) with an integer parameter N . In that case, it invokes the first version of E.PublicKeyGen (not the alternate LPR version presented above) with the specified value of N .

We remark that an analogous approach can be used to design an alternate encryption scheme based on LWE.

4.2. Key Switching

We start by reminding the reader that in the basic GLWE-based encryption scheme, the decryption equation for a ciphertext \mathbf{c} that encrypts m under key \mathbf{s} can be written as $m = [\Lambda_{\mathbf{c}}(\mathbf{s})]_q]_2$, where $\Lambda_{\mathbf{c}}(\mathbf{x})$ is a ciphertext-dependent linear equation over the coefficients of \mathbf{x} given by $\Lambda_{\mathbf{c}}(\mathbf{x}) = \langle \mathbf{c}, \mathbf{x} \rangle$.

Suppose now that we have two ciphertexts \mathbf{c}_1 and \mathbf{c}_2 encrypting m_1 and m_2 , respectively, under the same secret key \mathbf{s} . The way homomorphic multiplication is accomplished in Brakerski and Vaikuntanathan [2011b] is to consider the quadratic equation $Q_{\mathbf{c}_1, \mathbf{c}_2}(\mathbf{x}) := \Lambda_{\mathbf{c}_1}(\mathbf{x}) \cdot \Lambda_{\mathbf{c}_2}(\mathbf{x})$. Assuming the noise in the initial ciphertexts are small enough, we obtain

$$m_1 \cdot m_2 = [Q_{\mathbf{c}_1, \mathbf{c}_2}(\mathbf{s})]_q]_2,$$

as desired. If one wishes, one can view $Q_{\mathbf{c}_1, \mathbf{c}_2}(\mathbf{x})$ as a linear equation $\Lambda_{\mathbf{c}_1, \mathbf{c}_2}^{long}(\mathbf{x} \otimes \mathbf{x})$ over the coefficients of $\mathbf{x} \otimes \mathbf{x}$. Here, $\mathbf{x} \otimes \mathbf{x}$ is the tensoring of \mathbf{x} with itself, increasing the dimension to the square of the dimension of \mathbf{x} 's. Using this interpretation, the ciphertext represented by the coefficients of the linear equation L^{long} is decryptable by the long secret key $\mathbf{s} \otimes \mathbf{s}$ via the usual dot product. Of course, we cannot continue increasing the dimension like this indefinitely and preserve efficiency.

Thus, Brakerski and Vaikuntanathan convert the long ciphertext represented by the linear equation Λ^{long} and decryptable by the long tensored secret key $\mathbf{s} \otimes \mathbf{s}$ into a shorter ciphertext \mathbf{c}_2 that is decryptable by a different secret key \mathbf{s}_2 . (The secret keys need to be different to avoid a circular security issue, described in Section 2.3). Encryptions of $\mathbf{s}_1 \otimes \mathbf{s}_1$ under \mathbf{s}_2 are provided in the public key as a hint to facilitate this conversion. They call this the *relinearization* procedure.

The starting point of our key switching procedure is to observe that Brakerski and Vaikuntanathan's relinearization procedure is actually quite a bit more general. It can be used to not only reduce the dimension of the ciphertext, but more generally, can be used to transform a ciphertext \mathbf{c}_1 that is decryptable under one secret key vector \mathbf{s}_1

to a different ciphertext \mathbf{c}_2 that encrypts the same message, but is now decryptable under a second secret key vector \mathbf{s}_2 . The vectors $\mathbf{c}_2, \mathbf{s}_2$ may not necessarily be of lower degree or dimension than $\mathbf{c}_1, \mathbf{s}_1$. Because of this generality, we prefer to call this the *key switching* procedure, which we now describe in detail.

Two Useful Subroutines. The key switching procedure will use some subroutines that, given two vectors \mathbf{c} and \mathbf{s} , expand these vectors to get longer (higher-dimensional) vectors \mathbf{c}' and \mathbf{s}' such that $\langle \mathbf{c}', \mathbf{s}' \rangle = \langle \mathbf{c}, \mathbf{s} \rangle \bmod q$.

- BitDecomp($\mathbf{c} \in R_q^n, q$) decomposes \mathbf{c} into its bit representation. Namely, write $\mathbf{c} = \sum_{j=0}^{\lfloor \log q \rfloor} 2^j \cdot \mathbf{c}_j$ with all $\mathbf{c}_j \in R_2^n$. Output $(\mathbf{c}_0, \mathbf{c}_1, \dots, \mathbf{c}_{\lfloor \log q \rfloor}) \in R_2^{n \cdot \lceil \log q \rceil}$.
- Powersof2($\mathbf{s} \in R_q^n, q$) outputs $(\mathbf{s}, 2 \cdot \mathbf{s}, \dots, 2^{\lfloor \log q \rfloor} \cdot \mathbf{s}) \in R_q^{n \cdot \lceil \log q \rceil}$.

We observe the following elementary fact.

LEMMA 4.1. *For vectors \mathbf{c}, \mathbf{s} of equal length, we have*

$$\langle \text{BitDecomp}(\mathbf{c}, q), \text{Powersof2}(\mathbf{s}, q) \rangle = \langle \mathbf{c}, \mathbf{s} \rangle \bmod q .$$

PROOF. The proof follows by a simple calculation from the definitions of the previous procedures BitDecomp and Powersof2. In particular, we have

$$\begin{aligned} \langle \text{BitDecomp}(\mathbf{c}, q), \text{Powersof2}(\mathbf{s}, q) \rangle &= \sum_{j=0}^{\lfloor \log q \rfloor} \langle \mathbf{c}_j, 2^j \cdot \mathbf{s} \rangle = \sum_{j=0}^{\lfloor \log q \rfloor} \langle 2^j \cdot \mathbf{c}_j, \mathbf{s} \rangle \\ &= \left\langle \sum_{j=0}^{\lfloor \log q \rfloor} 2^j \cdot \mathbf{c}_j, \mathbf{s} \right\rangle = \langle \mathbf{c}, \mathbf{s} \rangle . \quad \square \end{aligned}$$

We remark that this obviously generalizes to decompositions with respect to bases other than the powers of 2. Also, as an optimization, if one knows a priori that \mathbf{c} has coefficients in $[0, B]$ for $B \ll q$, then BitDecomp can be optimized in the obvious way to output a shorter decomposition in $R_2^{n \cdot \lceil \log B \rceil}$.

The Key Switching Procedure. Key switching consists of two procedures: first, an algorithm $\text{SwitchKeyGen}(\mathbf{s}_1, \mathbf{s}_2, n_1, n_2, q)$, which takes as input the two secret key vectors, the respective dimensions of these vectors, and the modulus q , and outputs some auxiliary information $\tau_{\mathbf{s}_1 \rightarrow \mathbf{s}_2}$ that enables the switching; and second, an algorithm $\text{SwitchKey}(\tau_{\mathbf{s}_1 \rightarrow \mathbf{s}_2}, \mathbf{c}_1, n_1, n_2, q)$, that takes this auxiliary information and a ciphertext encrypted under \mathbf{s}_1 and outputs a new ciphertext \mathbf{c}_2 that encrypts the same message under the secret key \mathbf{s}_2 . (In the following, we often suppress the additional arguments n_1, n_2, q .)

$\text{SwitchKeyGen}(\mathbf{s}_1 \in R_q^{n_1}, \mathbf{s}_2 \in R_q^{n_2})$.

- Run $\mathbf{A} \leftarrow \mathbf{E}.\text{PublicKeyGen}(\mathbf{s}_2, N)$ for $N = n_1 \cdot \lceil \log q \rceil$.
- Add Powersof2($\mathbf{s}_1 \in R_q^{N_1}$) to \mathbf{A} 's first column to get a matrix \mathbf{B} . Output $\tau_{\mathbf{s}_1 \rightarrow \mathbf{s}_2} = \mathbf{B}$.

$\text{SwitchKey}(\tau_{\mathbf{s}_1 \rightarrow \mathbf{s}_2}, \mathbf{c}_1)$. Output $\mathbf{c}_2 = \text{BitDecomp}(\mathbf{c}_1)^T \cdot \mathbf{B} \in R_q^{n_2}$.

Note that, in SwitchKeyGen , the matrix \mathbf{A} basically consists of encryptions of 0 under the key \mathbf{s}_2 . Then, pieces of the key \mathbf{s}_1 are added to these encryptions of 0. Thus, in some sense, the matrix \mathbf{B} consists of encryptions of pieces of \mathbf{s}_1 (in a certain format) under the key \mathbf{s}_2 . We now establish that the key switching procedures are meaningful, in the sense that they preserve the correctness of decryption under the new key.

LEMMA 4.2. Let $\mathbf{s}_1, \mathbf{s}_2, q, \mathbf{A}, \mathbf{B} = \tau_{\mathbf{s}_1 \rightarrow \mathbf{s}_2}$ be as in $\text{SwitchKeyGen}(\mathbf{s}_1, \mathbf{s}_2)$, and let $\mathbf{A} \cdot \mathbf{s}_2 = 2\mathbf{e}_2 \in R_q^N$. Let $\mathbf{c}_1 \in R_q^{n_1}$ and $\mathbf{c}_2 \leftarrow \text{SwitchKey}(\tau_{\mathbf{s}_1 \rightarrow \mathbf{s}_2}, \mathbf{c}_1)$. Then,

$$\langle \mathbf{c}_2, \mathbf{s}_2 \rangle = 2 \langle \text{BitDecomp}(\mathbf{c}_1), \mathbf{e}_2 \rangle + \langle \mathbf{c}_1, \mathbf{s}_1 \rangle \bmod q.$$

PROOF. We prove this by the following sequence of equalities:

$$\begin{aligned} \langle \mathbf{c}_2, \mathbf{s}_2 \rangle &= \text{BitDecomp}(\mathbf{c}_1)^T \cdot \mathbf{B} \cdot \mathbf{s}_2 \\ &= \text{BitDecomp}(\mathbf{c}_1)^T \cdot (2\mathbf{e}_2 + \text{Powersof2}(\mathbf{s}_1)) \\ &= 2 \langle \text{BitDecomp}(\mathbf{c}_1), \mathbf{e}_2 \rangle + \langle \text{BitDecomp}(\mathbf{c}_1), \text{Powersof2}(\mathbf{s}_1) \rangle \\ &= 2 \langle \text{BitDecomp}(\mathbf{c}_1), \mathbf{e}_2 \rangle + \langle \mathbf{c}_1, \mathbf{s}_1 \rangle, \end{aligned}$$

where the first equality is by the definition of SwitchKey , the second by the definition of SwitchKeyGen , and the last by Lemma 4.1. \square

Note that the dot product of $\text{BitDecomp}(\mathbf{c}_1)$ and \mathbf{e}_2 is small, since $\text{BitDecomp}(\mathbf{c}_1)$ is in R_2^N . Overall, the consequence of this lemma is that \mathbf{c}_2 is a valid encryption of m under key \mathbf{s}_2 , with noise that is larger by a small additive factor.

We will need in the sequel the following lemma which says that the key switching parameter generated by $\text{SwitchKeyGen}(\mathbf{s}_1, \mathbf{s}_2)$ is computationally indistinguishable from random for any $\mathbf{s}_1 \in R_q^{n_1}$ and a uniformly random $\mathbf{s}_2 \in R_q^{n_2}$. This is simply because of the properties of the basic scheme E , which states that the matrix $\mathbf{A} \leftarrow E.\text{PublicKeyGen}(\mathbf{s}_2, N)$ is computationally indistinguishable from uniform.

LEMMA 4.3 (SECURITY). For every $\mathbf{s}_1 \in R_q^{n_2}$ and a uniformly random $\mathbf{s}_2 \leftarrow R_q^{n_2}$, the following two distributions are computationally indistinguishable (even to a distinguisher that knows \mathbf{s}_1):

$$\begin{aligned} \left\{ (\mathbf{A}, \tau_{\mathbf{s}_1 \rightarrow \mathbf{s}_2}) : \mathbf{A} \leftarrow E.\text{PublicKeyGen}(\mathbf{s}_2, N), \tau_{\mathbf{s}_1 \rightarrow \mathbf{s}_2} \leftarrow \text{SwitchKeyGen}(\mathbf{s}_1, \mathbf{s}_2) \right\} &\approx \\ \left\{ (\mathbf{A}, \tau_{\mathbf{s}_1 \rightarrow \mathbf{s}_2}) : \mathbf{A} \leftarrow R_q^{N \times (n_2+1)}, \tau_{\mathbf{s}_1 \rightarrow \mathbf{s}_2} \leftarrow R_q^{N \times (n_2+1)} \right\}. \end{aligned}$$

Here, the randomness is over the choice of $\mathbf{s}_2 \leftarrow R_q^{n_2}$, and the coins of $E.\text{PublicKeyGen}$ and SwitchKeyGen .

As a final remark in this section, we note that in the follow-up work of Gentry et al. [2012c], which aimed at a practical implementation of this work, the authors chose to not use bit decomposition in the key switching process. Instead, they control the noise in a different method whose asymptotic performance may be worse, but which turns out to perform better in practice.

4.3. Modulus Switching

Suppose \mathbf{c} is a valid encryption of m under \mathbf{s} modulo q (i.e., $m = [\langle \mathbf{c}, \mathbf{s} \rangle]_q]_2$), and that \mathbf{s} is a *short* vector. Suppose also that \mathbf{c}' is basically a simple scaling of \mathbf{c} —in particular, \mathbf{c}' is the R -vector closest to $(p/q) \cdot \mathbf{c}$ such that $\mathbf{c}' = \mathbf{c} \bmod 2$. Then, it turns out (subject to some qualifications) that \mathbf{c}' is a valid encryption of m under \mathbf{s} modulo p using the usual decryption equation, that is, $m = [\langle \mathbf{c}', \mathbf{s} \rangle]_p]_2$! In other words, we can change the inner modulus in the decryption equation (e.g., to a smaller number) while preserving the correctness of decryption under the same secret key. The essence of this modulus switching idea, a variant of Brakerski and Vaikuntanathan's modulus reduction technique, is formally captured in Lemma 4.6.

Definition 4.4 (Scale). For integer vector \mathbf{x} and integers $q > p > m$, we define $\mathbf{x}' \leftarrow \text{Scale}(\mathbf{x}, q, p, r)$ to be the R -vector closest to $(p/q) \cdot \mathbf{x}$ that satisfies $\mathbf{x}' = \mathbf{x} \bmod r$.

Definition 4.5 ($\ell_1^{(R)}$ norm). The (usual) norm $\ell_1(\mathbf{s})$ over the reals equals $\sum_i \|\mathbf{s}[i]\|$. We extend this to our ring R as follows: $\ell_1^{(R)}(\mathbf{s})$ for $\mathbf{s} \in R^n$ is defined as $\sum_i \|\mathbf{s}[i]\|$.

LEMMA 4.6. Let d be the degree of the ring (e.g., $d = 1$ when $R = \mathbb{Z}$). Let $q > p > r$ be positive integers satisfying $q = p = 1 \bmod r$. Let $\mathbf{c} \in R^n$ and $\mathbf{c}' \leftarrow \text{Scale}(\mathbf{c}, q, p, r)$. Then, for any $\mathbf{s} \in R^n$ with $\|\langle \mathbf{c}, \mathbf{s} \rangle\|_q < q/2 - (q/p) \cdot \gamma_R \cdot (r/2) \cdot \sqrt{d} \cdot \ell_1^{(R)}(\mathbf{s})$, we have

$$\begin{aligned} \lfloor \langle \mathbf{c}', \mathbf{s} \rangle \rfloor_p &= \lfloor \langle \mathbf{c}, \mathbf{s} \rangle \rfloor_q \bmod r \quad \text{and} \\ \|\lfloor \langle \mathbf{c}', \mathbf{s} \rangle \rfloor_p\| &< (p/q) \cdot \|\lfloor \langle \mathbf{c}, \mathbf{s} \rangle \rfloor_q\| + \gamma_R \cdot (r/2) \cdot \sqrt{d} \cdot \ell_1^{(R)}(\mathbf{s}). \end{aligned}$$

PROOF (LEMMA 4.6). We have

$$\lfloor \langle \mathbf{c}, \mathbf{s} \rangle \rfloor_q = \langle \mathbf{c}, \mathbf{s} \rangle - kq$$

for some $k \in R$. For the same k , let

$$e_p = \langle \mathbf{c}', \mathbf{s} \rangle - kp \in R.$$

Note that $e_p = \lfloor \langle \mathbf{c}', \mathbf{s} \rangle \rfloor_p \bmod p$. We claim that $\|e_p\|$ is so small that $e_p = \lfloor \langle \mathbf{c}', \mathbf{s} \rangle \rfloor_p$. We have

$$\begin{aligned} \|e_p\| &= \| -kp + \langle (p/q) \cdot \mathbf{c}, \mathbf{s} \rangle + \langle \mathbf{c}' - (p/q) \cdot \mathbf{c}, \mathbf{s} \rangle \| \\ &\leq \| -kp + \langle (p/q) \cdot \mathbf{c}, \mathbf{s} \rangle \| + \| \langle \mathbf{c}' - (p/q) \cdot \mathbf{c}, \mathbf{s} \rangle \| \\ &\leq (p/q) \cdot \|\lfloor \langle \mathbf{c}, \mathbf{s} \rangle \rfloor_q\| + \gamma_R \cdot \sum_{j=1}^n \|\mathbf{c}'[j] - (p/q) \cdot \mathbf{c}[j]\| \cdot \|\mathbf{s}[j]\| \\ &\leq (p/q) \cdot \|\lfloor \langle \mathbf{c}, \mathbf{s} \rangle \rfloor_q\| + \gamma_R \cdot (r/2) \cdot \sqrt{d} \cdot \ell_1^{(R)}(\mathbf{s}) \\ &< p/2. \end{aligned}$$

Furthermore, modulo r , we have $\lfloor \langle \mathbf{c}', \mathbf{s} \rangle \rfloor_p = e_p = \langle \mathbf{c}', \mathbf{s} \rangle - kp = \langle \mathbf{c}, \mathbf{s} \rangle - kq = \lfloor \langle \mathbf{c}, \mathbf{s} \rangle \rfloor_q$. \square

The lemma implies that an evaluator, who does not know the secret key but instead only knows a bound on its length, could potentially transform a ciphertext \mathbf{c} that encrypts m under key \mathbf{s} for modulus q , that is, $m = \lfloor \langle \mathbf{c}, \mathbf{s} \rangle \rfloor_q \bmod r$, into a ciphertext \mathbf{c}' that encrypts m under the same key \mathbf{s} for modulus p , that is, $m = \lfloor \langle \mathbf{c}', \mathbf{s} \rangle \rfloor_p \bmod r$. Specifically, the following corollary follows immediately from Lemma 4.6.

COROLLARY 4.7. Let p and q be two odd moduli. Suppose \mathbf{c} is an encryption of bit m under key \mathbf{s} for modulus q , that is, $m = \lfloor \langle \mathbf{c}, \mathbf{s} \rangle \rfloor_q \bmod r$. Moreover, suppose that \mathbf{s} is a fairly short key and the noise $e_q \leftarrow \lfloor \langle \mathbf{c}, \mathbf{s} \rangle \rfloor_q$ has small magnitude—precisely, assume that $\|e_q\| < q/2 - (q/p) \cdot (r/2) \cdot \sqrt{d} \cdot \gamma_R \cdot \ell_1^{(R)}(\mathbf{s})$. Then $\mathbf{c}' \leftarrow \text{Scale}(\mathbf{c}, q, p, r)$ is an encryption of bit m under key \mathbf{s} for modulus p , that is, $m = \lfloor \langle \mathbf{c}', \mathbf{s} \rangle \rfloor_p \bmod r$. The noise $e_p = \lfloor \langle \mathbf{c}', \mathbf{s} \rangle \rfloor_p$ of the new ciphertext has magnitude at most $(p/q) \cdot \|\lfloor \langle \mathbf{c}, \mathbf{s} \rangle \rfloor_q\| + \gamma_R \cdot (r/2) \cdot \sqrt{d} \cdot \ell_1^{(R)}(\mathbf{s})$.

Amazingly, assuming p is smaller than q and \mathbf{s} has coefficients that are small in relation to q , this trick permits the evaluator to reduce the magnitude of the noise without knowing the secret key! (Of course, this is also what Gentry's bootstrapping transformation accomplishes, but in a much more complicated way.)

4.4. (Leveled) FHE Based without Bootstrapping

We now present our leveled fully homomorphic encryption scheme. Given the machinery that we have described in the previous subsections, the scheme itself is remarkably simple.

In our scheme, we will use a parameter L indicating the number of levels of arithmetic circuit that we want our FHE scheme to be capable of evaluating. Note that this is an exponential improvement over prior schemes that would typically use a parameter d indicating the *degree* of the polynomials to be evaluated.

- $\text{FHE}.\text{Setup}(1^\lambda, 1^L)$. Takes as input the security parameter and a number of levels L . Let $\mu = \mu(\lambda, L, b) = \theta(\log \lambda + \log L)$ be a parameter that we will specify in detail later. For $j = L$ (input level of circuit) to 0 (output level), run $\text{params}_j \leftarrow \text{E}.\text{Setup}(1^\lambda, 1^{(j+1)\cdot\mu}, b)$ to obtain a ladder of parameters, including a ladder of decreasing moduli from q_L ($(L+1) \cdot \mu$ bits) down to q_0 (μ bits). (The ring degree d_j , dimension n_j , and noise distribution χ_j do not necessarily need to vary (decrease) with the circuit level. In the following procedure, we allow n_j and χ_j to vary but defer the case of decreasing d_j to Section 3.)
- $\text{FHE}.\text{KeyGen}(\{\text{params}_j\})$. For $j = L$ down to 0, do the following.
 - (1) Run $\mathbf{s}_j \leftarrow \text{E}.\text{SecretKeyGen}(\text{params}_j)$ and $\mathbf{A}_j \leftarrow \text{E}.\text{PublicKeyGen}(\text{params}_j, \mathbf{s}_j)$.
 - (2) Set $\mathbf{s}'_j \leftarrow \mathbf{s}_j \otimes \mathbf{s}_j \in R_{q_j}^{\binom{n_j+1}{2}}$, that is, \mathbf{s}'_j is a tensoring of \mathbf{s}_j with itself whose coefficients are each the product of two coefficients of \mathbf{s}_j in R_{q_j} .
 - (3) Run $\tau_{\mathbf{s}'_{j+1} \rightarrow \mathbf{s}_j} \leftarrow \text{SwitchKeyGen}(\mathbf{s}'_{j+1}, \mathbf{s}_j)$. (Omit this step when $j = L$.)
 The secret key sk consists of the \mathbf{s}_j 's and the public key pk consists of the \mathbf{A}_j 's and $\tau_{\mathbf{s}'_{j+1} \rightarrow \mathbf{s}_j}$'s.
- $\text{FHE}.\text{Enc}(\text{params}, pk, m)$. Take a message in R_2 . Run $\text{E}.\text{Enc}(\text{params}_L, \mathbf{A}_L, m)$.
- $\text{FHE}.\text{Dec}(\text{params}, sk, \mathbf{c})$. Suppose the ciphertext is under key \mathbf{s}_j . Run $\text{E}.\text{Dec}(\text{params}_j, \mathbf{s}_j, \mathbf{c})$. (The ciphertext could be augmented with an index indicating which level it belongs to.)
- $\text{FHE}.\text{Eval}(pk, f, \mathbf{c}_1, \dots, \mathbf{c}_\ell)$. Take as input a circuit f with ℓ inputs, as well as ciphertexts $\mathbf{c}_1, \dots, \mathbf{c}_\ell$. We will assume without loss of generality that f is a leveled circuit composed of layers of alternating addition and multiplication gates. $\text{FHE}.\text{Eval}$ will invoke $\text{FHE}.\text{Add}$ and $\text{FHE}.\text{Mult}$ described next to compute the circuit layer by layer, starting from the inputs. In addition, after each multiplication layer, it will invoke the ciphertext refreshing procedure $\text{FHE}.\text{Refresh}$ to reduce the noise in the ciphertexts and move it to a different level. These procedures are described in detail next.
- $\text{FHE}.\text{Add}(pk, \mathbf{c}_1, \mathbf{c}_2)$. Takes two ciphertexts encrypted under the same \mathbf{s}_j . (If needed, use $\text{FHE}.\text{Refresh}$ to make it so.) Set $\mathbf{c}_3 \leftarrow \mathbf{c}_1 + \mathbf{c}_2 \bmod q_j$. Interpret \mathbf{c}_3 as a ciphertext under \mathbf{s}'_j (\mathbf{s}'_j 's coefficients include all of \mathbf{s}_j 's, since $\mathbf{s}'_j = \mathbf{s}_j \otimes \mathbf{s}_j$ and \mathbf{s}_j 's first coefficient is (1) and output

$$\mathbf{c}_4 \leftarrow \text{FHE}.\text{Refresh}(\mathbf{c}_3, \tau_{\mathbf{s}'_j \rightarrow \mathbf{s}_{j-1}}, q_j, q_{j-1}).$$

- $\text{FHE}.\text{Mult}(pk, \mathbf{c}_1, \mathbf{c}_2)$. Takes two ciphertexts encrypted under the same \mathbf{s}_j . (If needed, use $\text{FHE}.\text{Refresh}$ to make it so.) First, multiply: the new ciphertext, under the secret key $\mathbf{s}'_j = \mathbf{s}_j \otimes \mathbf{s}_j$, is the coefficient vector \mathbf{c}_3 of the linear equation

$$\Lambda_{\mathbf{c}_1, \mathbf{c}_2}^{\text{long}}(\mathbf{x} \otimes \mathbf{x}) := \langle \mathbf{c}_1 \otimes \mathbf{c}_2, \mathbf{x} \otimes \mathbf{x} \rangle.$$

Then, output

$$\mathbf{c}_4 \leftarrow \text{FHE}.\text{Refresh}(\mathbf{c}_3, \tau_{\mathbf{s}'_j \rightarrow \mathbf{s}_{j-1}}, q_j, q_{j-1}).$$

— $\text{FHE}.\text{Refresh}(\mathbf{c}, \tau_{\mathbf{s}'_j \rightarrow \mathbf{s}_{j-1}}, q_j, q_{j-1})$. Takes a ciphertext encrypted under \mathbf{s}'_j , the auxiliary information $\tau_{\mathbf{s}'_j \rightarrow \mathbf{s}_{j-1}}$ to facilitate key switching, and the current and next moduli q_j and q_{j-1} . Do the following.

(1) *Switch Keys*. Set $\mathbf{c}_1 \leftarrow \text{SwitchKey}(\tau_{\mathbf{s}'_j \rightarrow \mathbf{s}_{j-1}}, \mathbf{c}, q_j)$, a ciphertext under the key \mathbf{s}_{j-1} for modulus q_j .

(2) *Switch Moduli*. Set $\mathbf{c}_2 \leftarrow \text{Scale}(\mathbf{c}_1, q_j, q_{j-1}, 2)$, a ciphertext under the key \mathbf{s}_{j-1} for modulus q_{j-1} .

Remark 4.8. We mention the obvious fact that since addition increases the noise much more slowly than multiplication, one does not necessarily need to refresh after additions, even high fan-in ones.

The key step of our new FHE scheme is the Refresh procedure. If the modulus q_{j-1} is chosen to be smaller than q_j by a sufficient multiplicative factor, then Corollary 4.7 implies that the noise of the ciphertext output by Refresh is smaller than that of the input ciphertext, that is, the ciphertext will indeed be a “refreshed” encryption of the same value. We elaborate on this analysis in the next section.

One can reasonably argue that this scheme is not FHE without bootstrapping, since $\tau_{\mathbf{s}'_j \rightarrow \mathbf{s}_{j-1}}$ can be viewed as an encrypted secret key, and the SwitchKey step can be viewed as a homomorphic evaluation of the decryption function. We prefer not to view the SwitchKey step this way. While there is some high-level resemblance, the low-level details are very different, a difference that becomes tangible in the much better asymptotic performance. To the extent that it performs decryption, SwitchKey does so very efficiently using an efficient (not bitwise) representation of the secret key that allows this step to be computed in quasilinear time for the RLWE instantiation, below the quadratic lower bound for bootstrapping. Certainly SwitchKey does not use the usual ponderous approach of representing the decryption function as a boolean circuit to be traversed homomorphically. Another difference is that the SwitchKey step does not actually reduce the noise level (as bootstrapping does); rather, the noise is reduced by the Scale step.

5. PARAMETER SETTINGS: CORRECTNESS, PERFORMANCE, AND SECURITY

Here, we will show how to set the parameters of our FHE scheme to get the desired homomorphic capacity. Mostly, this involves analyzing each of the steps within $\text{FHE}.\text{Add}$ and $\text{FHE}.\text{Mult}$ —namely, the addition or multiplication itself, and then the SwitchKey and Scale steps that make up $\text{FHE}.\text{Refresh}$ —to establish that the output of each step is a decryptable ciphertext with bounded noise. This analysis will lead to concrete suggestions for how to set the ladder of moduli and to asymptotic bounds on the performance of the scheme.

Let us begin by considering how much noise $\text{FHE}.\text{Enc}$ introduces initially. Throughout, B_χ denotes a bound such that R -elements sampled from the noise distribution χ have length at most B_χ (with probability 1).

5.1. The Initial Noise from $\text{FHE}.\text{Enc}$

Recall that $\text{FHE}.\text{Enc}$ simply invokes the encryption algorithm for the underlying public-key encryption scheme, namely, $\text{E}.\text{Enc}$, for suitable parameters params_L that depend on λ and L . In turn, the noise of ciphertexts output by $\text{E}.\text{Enc}$ depends on the noise of the initial ciphertext(s) (the encryption(s) of 0) implicit in the matrix \mathbf{A} output by $\text{E}.\text{PublicKeyGen}$, whose noise distribution is dictated by the distribution χ .

LEMMA 5.1. *Let q, d, n, N be the parameters associated to $\text{FHE}.\text{Enc}$. Let γ_R be the expansion factor associated to R . (γ_R and d are both 1 in the LWE case $R = \mathbb{Z}$.) The*

Euclidean norm of the noise in ciphertexts output by FHE.Enc is at most $\sqrt{d} + 2 \cdot \gamma_R \cdot \sqrt{d} \cdot N \cdot B_\chi$.

PROOF. We have $\mathbf{A} \cdot \mathbf{s} = 2\mathbf{e}$, where $\mathbf{s} \leftarrow \text{E.SecretKeyGen}$, $\mathbf{A} \leftarrow \text{E.PublicKeyGen}(\mathbf{s}, N)$, and $\mathbf{e} \leftarrow \chi^N$. Recall that encryption works as follows: $\mathbf{c} \leftarrow \mathbf{m} + \mathbf{A}^T \mathbf{r} \bmod q$, where $\mathbf{r} \in R_2^N$. We have that the noise of this ciphertext is $[\langle \mathbf{c}, \mathbf{s} \rangle]_q = [m + 2\langle \mathbf{r}, \mathbf{e} \rangle]_q$. The magnitude of this element is at most

$$\sqrt{d} + 2 \cdot \gamma_R \cdot \sum_{j=1}^N \|\mathbf{r}[j]\| \cdot \|\mathbf{e}[j]\| \leq \sqrt{d} + 2 \cdot \gamma_R \cdot \sqrt{d} \cdot N \cdot B_\chi. \quad \square$$

One can easily obtain a similar small bound on the noise of ciphertexts output by LPR encryption in the RLWE setting: a small polynomial in the security parameter λ , L , and $\log q$.

The correctness of decryption for ciphertexts output by FHE.Enc, assuming the noise bound is less than $q/2$, follows directly from the correctness of the basic encryption and decryption algorithms E.Enc and E.Dec.

5.2. Correctness and Performance of FHE.Add and FHE.Mult (before FHE.Refresh)

Consider FHE.Mult. One begins FHE.Mult($pk, \mathbf{c}_1, \mathbf{c}_2$) with two ciphertexts under key \mathbf{s}_j for modulus q_j that have noises $e_i = [L_{\mathbf{c}_i}(\mathbf{s}_j)]_{q_j}$, where $L_{\mathbf{c}_i}(\mathbf{x})$ is simply the dot product $\langle \mathbf{c}_i, \mathbf{x} \rangle$. To multiply together two ciphertexts, one multiplies together these two linear equations to obtain a quadratic equation $Q_{\mathbf{c}_1, \mathbf{c}_2}(\mathbf{x}) \leftarrow L_{\mathbf{c}_1}(\mathbf{x}) \cdot L_{\mathbf{c}_2}(\mathbf{x})$, and then interprets this quadratic equation as a linear equation $L_{\mathbf{c}_1, \mathbf{c}_2}^{\text{long}}(\mathbf{x} \otimes \mathbf{x}) = Q_{\mathbf{c}_1, \mathbf{c}_2}(\mathbf{x})$ over the tensored vector $\mathbf{x} \otimes \mathbf{x}$. The coefficients of this long linear equation compose the new ciphertext vector \mathbf{c}_3 . Clearly, $[\langle \mathbf{c}_3, \mathbf{s}_j \otimes \mathbf{s}_j \rangle]_{q_j} = [L_{\mathbf{c}_1, \mathbf{c}_2}^{\text{long}}(\mathbf{s}_j \otimes \mathbf{s}_j)]_{q_j} = [e_1 \cdot e_2]_{q_j}$. Thus, if the noises of \mathbf{c}_1 and \mathbf{c}_2 have norm at most B , then the noise of \mathbf{c}_3 has norm at most $\gamma_R \cdot B^2$, where γ_R is the expansion factor of R . If this norm is less than $q_j/2$, then decryption works correctly. In particular, if $m_i = [\langle \mathbf{c}_i, \mathbf{s}_j \rangle]_{q_j} \|_2 = [e_i]_2$ for $i \in \{1, 2\}$, then over R_2 we have $[\langle \mathbf{c}_3, \mathbf{s}_j \otimes \mathbf{s}_j \rangle]_{q_j} \|_2 = [[e_1 \cdot e_2]_{q_j}]_2 = [e_1 \cdot e_2]_2 = [e_1]_2 \cdot [e_2]_2 = m_1 \cdot m_2$. That is, correctness is preserved as long as this noise does not wrap modulo q_j .

The correctness of FHE.Add and FHE.Mult, before the FHE.Refresh step is performed, is formally captured in the following lemmas.

LEMMA 5.2. Let \mathbf{c}_1 and \mathbf{c}_2 be two ciphertexts under key \mathbf{s}_j for modulus q_j , where $\|[\langle \mathbf{c}_i, \mathbf{s}_j \rangle]_{q_j}\| \leq B$ and $m_i = [[\langle \mathbf{c}_i, \mathbf{s}_j \rangle]_{q_j}]_2$. Let $\mathbf{s}'_j = \mathbf{s}_j \otimes \mathbf{s}_j$, where the nonquadratic coefficients of \mathbf{s}'_j (namely, the '1' and the coefficients of \mathbf{s}_j) are placed first. Let $\mathbf{c}' = \mathbf{c}_1 + \mathbf{c}_2$, and pad \mathbf{c}' with zeros to get a vector \mathbf{c}_3 such that $\langle \mathbf{c}_3, \mathbf{s}'_j \rangle = \langle \mathbf{c}', \mathbf{s}_j \rangle$. The noise $[\langle \mathbf{c}_3, \mathbf{s}'_j \rangle]_{q_j}$ has norm at most $2B$. If $2B < q_j/2$, \mathbf{c}_3 is an encryption of $m_1 + m_2$ under key \mathbf{s}'_j for modulus q_j , that is, $m_1 \cdot m_2 = [[\langle \mathbf{c}_3, \mathbf{s}'_j \rangle]_{q_j}]_2$.

LEMMA 5.3. Let \mathbf{c}_1 and \mathbf{c}_2 be two ciphertexts under key \mathbf{s}_j for modulus q_j , where $\|[\langle \mathbf{c}_i, \mathbf{s}_j \rangle]_{q_j}\| \leq B$ and $m_i = [[\langle \mathbf{c}_i, \mathbf{s}_j \rangle]_{q_j}]_2$. Let the linear equation $L_{\mathbf{c}_1, \mathbf{c}_2}^{\text{long}}(\mathbf{x} \otimes \mathbf{x})$ be as previously defined, let \mathbf{c}_3 be the coefficient vector of this linear equation, and let $\mathbf{s}'_j = \mathbf{s}_j \otimes \mathbf{s}_j$. The noise $[\langle \mathbf{c}_3, \mathbf{s}'_j \rangle]_{q_j}$ has norm at most $\gamma_R \cdot B^2$. If $\gamma_R \cdot B^2 < q_j/2$, \mathbf{c}_3 is an encryption of $m_1 \cdot m_2$ under key \mathbf{s}'_j for modulus q_j , that is, $m_1 \cdot m_2 = [[\langle \mathbf{c}_3, \mathbf{s}'_j \rangle]_{q_j}]_2$.

The computation needed to compute the tensored ciphertext c_3 is $\tilde{O}(d_j n_j^2 \log q_j)$. For the RLWE instantiation, since $n_j = 1$ and since (as we will see) d_j (resp. $\log q_j$) depend only quasilinearly (resp. logarithmically) on the security parameter and linearly (resp. linearly) on L , the computation here is only quasilinear in the security parameter. For the LWE instantiation, the computation is quasiquadratic.

5.3. Correctness and Performance of FHE.Refresh

FHE.Refresh consists of two steps: Switch Keys and Switch Moduli. We address each of these steps in turn.

Correctness and Performance of the Switch-Key Step. In the Switch-Keys step, we take as input a ciphertext \mathbf{c} under key \mathbf{s}'_j for modulus q_j and set $\mathbf{c}_1 \leftarrow \text{SwitchKey}(\tau_{\mathbf{s}'_j \rightarrow \mathbf{s}_{j-1}}, \mathbf{c}, q_j)$, a ciphertext under the key \mathbf{s}_{j-1} for modulus q_j . In Lemma 4.2, we proved the correctness of key switching and showed that the noise grows by the additive factor $2 \langle \text{BitDecomp}(\mathbf{c}, q_j), \mathbf{e} \rangle$, where $\text{BitDecomp}(\mathbf{c}, q_j)$ is a (short) bit-vector and \mathbf{e} is a (short and fresh) noise vector with elements sampled from χ . In particular, if the noise originally had norm B , then after the Switch-Keys step, it has norm at most $B + 2 \cdot \gamma_R \cdot B_\chi \cdot \sum_{i=1}^{w_j} \|\text{BitDecomp}(\mathbf{c}, q_j)[i]\| \leq B + 2 \cdot \gamma_R \cdot B_\chi \cdot w_j \cdot \sqrt{d_j}$, where $w_j \leq \binom{n_j+1}{2} \cdot \lceil \log q_j \rceil$ is the dimension of $\text{BitDecomp}(\mathbf{c}, q_j)$.

We capture the correctness of the Switch-Key step in the following lemma.

LEMMA 5.4. *Let \mathbf{c} be a ciphertext under the key $\mathbf{s}'_j = \mathbf{s}_j \otimes \mathbf{s}_j$ for modulus q_j such that $e_1 \leftarrow [\langle \mathbf{c}, \mathbf{s}'_j \rangle]_{q_j}$ has norm at most B and $m = [e_1]_2$. Let $\mathbf{c}_1 \leftarrow \text{SwitchKey}(\tau_{\mathbf{s}'_j \rightarrow \mathbf{s}_{j-1}}, \mathbf{c}, q_j)$, and let $e_2 = [\langle \mathbf{c}_1, \mathbf{s}_{j-1} \rangle]_{q_j}$. Then, e_2 (the new noise) has norm at most $B + 2 \cdot \gamma_R \cdot B_\chi \cdot \binom{n_j+1}{2} \cdot \lceil \log q_j \rceil \cdot \sqrt{d_j}$ and (assuming this noise norm is less than $q_j/2$) we have $m = [e_2]_2$.*

The Switch-Key step involves multiplying the transpose of w_j -dimensional vector $\text{BitDecomp}(\mathbf{c}, q_j)$ with a $w_j \times (n_j + 1)$ matrix \mathbf{B} . This computation is $\tilde{O}(d_j n_j^3 \log^2 q_j)$. Still this is quasilinear in the RLWE instantiation.

Correctness and Performance of the Switch-Moduli Step. The Switch-Moduli step takes as input a ciphertext \mathbf{c}_1 under the secret bit-vector \mathbf{s}_{j-1} for the modulus q_j , and outputs the ciphertext $\mathbf{c}_2 \leftarrow \text{Scale}(\mathbf{c}_1, q_j, q_{j-1}, 2)$, which we claim to be a ciphertext under key \mathbf{s}_{j-1} for modulus q_{j-1} . Note that \mathbf{s}_{j-1} is a *short* secret key. By Corollary 4.7 and using the fact that $\ell_1(\mathbf{s}_{j-1}) \leq (n_{j-1} + 1) \cdot B_\chi$, the following is true: if the noise of \mathbf{c}_1 has norm at most $B < q_j/2 - (q_j/q_{j-1}) \cdot \sqrt{d_j} \cdot \gamma_R \cdot (n_{j-1} + 1) \cdot B_\chi$, then correctness is preserved and the noise of \mathbf{c}_2 is bounded by $(q_{j-1}/q_j) \cdot B + \sqrt{d_j} \cdot \gamma_R \cdot (n_{j-1} + 1) \cdot B_\chi$. Of course, the key feature of this step for our purposes is that switching moduli may reduce the size of the moduli when $q_{j-1} < q_j$.

We capture the correctness of the Switch-Moduli step in the following lemma.

LEMMA 5.5. *Let \mathbf{c}_1 be a ciphertext under the key \mathbf{s}_{j-1} , sampled from $\chi^{n_{j-1}}$, such that $e_j \leftarrow [\langle \mathbf{c}_1, \mathbf{s}_{j-1} \rangle]_{q_j}$ has norm at most B and $m = [e_j]_2$. Let $\mathbf{c}_2 \leftarrow \text{Scale}(\mathbf{c}_1, q_j, q_{j-1}, 2)$, and let $e_{j-1} = [\langle \mathbf{c}_2, \mathbf{s}_{j-1} \rangle]_{q_{j-1}}$. Then, e_{j-1} (the new noise) has norm at most $(q_{j-1}/q_j) \cdot B + \sqrt{d_j} \cdot \gamma_R \cdot (n_{j-1} + 1) \cdot B_\chi$, and (assuming this noise has norm less than $q_{j-1}/2$) we have $m = [e_{j-1}]_2$.*

The computation in the Switch-Moduli step is $\tilde{O}(d_j n_{j-1} \log q_j)$.

5.4. Putting the Pieces Together: Parameters, Correctness, Performance

So far we have established that the scheme is correct, assuming that the noise does not wrap modulo q_j or q_{j-1} . Now we need to show that we can set the parameters of the scheme to ensure that such wrapping never occurs.

Our strategy for setting the parameters is to pick a universal bound B on the noise length, and then prove, for all j , that a valid ciphertext under key \mathbf{s}_j for modulus q_j has noise length at most B . This bound B is quite small: polynomial in λ and $\log q_L$, where q_L is the largest modulus in our ladder. It is clear that such a bound B holds for fresh ciphertexts output by FHE.Enc. (Recall the discussion from Section 4.1, where we explained that we use a noise distribution χ that is essentially independent of the modulus.) The remainder of the proof is by induction, that is, we will show that if the bound holds for two ciphertexts $\mathbf{c}_1, \mathbf{c}_2$ at level j , our lemmas imply that the bound also holds for the ciphertext $\mathbf{c}' \leftarrow \text{FHE.Mult}(pk, \mathbf{c}_1, \mathbf{c}_2)$ at level $j - 1$. (FHE.Mult increases the noise strictly more in the worst case than FHE.Add for any reasonable choice of parameters.)

Specifically, after the first step of FHE.Mult (without the Refresh step), the noise has length at most $\gamma_R \cdot B^2$. Then, we apply the SwitchKey function, which introduces an additive term $\eta_{\text{SwitchKey},j}$. Finally, we apply the Scale function. The noise is now at most

$$(q_{j-1}/q_j) \cdot (\gamma_R \cdot B^2 + \eta_{\text{SwitchKey},j}) + \eta_{\text{Scale},j},$$

where $\eta_{\text{Scale},j}$ is another additive term. Now we want to choose our parameters so that this bound is at most B .

Suppose we set our ladder of moduli and the bound B such that the following two properties hold.

- Property 1. $B \geq 2 \cdot (\eta_{\text{Scale},j} + \eta_{\text{SwitchKey},j})$ for all j .
- Property 2. $q_j/q_{j-1} \geq 2 \cdot B \cdot \gamma_R$ for all j .

Then we have

$$\begin{aligned} & (q_{j-1}/q_j) \cdot (\gamma_R \cdot B^2 + \eta_{\text{SwitchKey},j}) + \eta_{\text{Scale},j} \\ & < (q_{j-1}/q_j) \cdot \gamma_R \cdot B^2 + \eta_{\text{Scale},j} + \eta_{\text{SwitchKey},j} \\ & \leq \frac{1}{2 \cdot B \cdot \gamma_R} \cdot \gamma_R \cdot B^2 + \frac{1}{2} \cdot B \\ & \leq B. \end{aligned}$$

It only remains to set our ladder of moduli and B so that Properties 1 and 2 hold.

Unfortunately, there is some circularity in Properties 1 and 2: q_L depends on B , which depends on q_L , albeit only polylogarithmically. However, it is easy to see that this circularity is not fatal, for example, by setting $B = \tilde{O}(B_\chi n_L^2 \sqrt{d_L} \gamma_R L) = \text{poly}(\lambda) \cdot \tilde{O}(L)$ and $q_j = (2B\gamma_R)^{j+1}$. We can therefore set q_j as a modulus having $(j+1) \cdot \mu$ bits for some $\mu = \theta(\log \lambda + \log L)$ with small hidden constant.

Overall, we have that q_L , the largest modulus used in the system, is $\theta(L \cdot (\log \lambda + \log L))$ bits, and $d_L \cdot n_L$ must be approximately that number times λ for 2^λ security.

The dependence of the complexity on L emerges mostly from its dependence on $\log(q)$. The dominant term is quadratic, coming from the complexity of the switch-key operation.

THEOREM 5.6. *For some $\mu = \theta(\log \lambda + \log L)$, FHE is a correct L -leveled FHE scheme: specifically, it correctly evaluates circuits of depth L with Add and Mult gates over R_2 . The per-gate computation is $\tilde{O}(d_L \cdot n_L^3 \cdot \log^2 q_j) = \tilde{O}(d_L \cdot n_L^3 \cdot L^2)$. For the LWE case (where*

$d = 1$), the per-gate computation is $\tilde{O}(\lambda^3 \cdot L^5)$. For the RLWE case (where $n = 1$), the per-gate computation is $\tilde{O}(\lambda \cdot L^3)$.

The bottom line is that we have a RLWE-based leveled FHE scheme with per-gate computation that is only quasilinear in the security parameter, albeit with somewhat high dependence on the number of levels in the circuit.

Let us pause at this point to reconsider the performance of previous FHE schemes in comparison to our new scheme. Specifically, as we discussed in the Introduction, in previous SWHE schemes, the ciphertext size is at least $\tilde{O}(\lambda \cdot d^2)$, where d is the *degree* of the circuit being evaluated. One may view our new scheme as a very powerful SWHE scheme in which this dependence on degree has been replaced with a similar dependence on *depth*. (Recall the degree of a circuit may be exponential in its depth.) Since polynomial-size circuits have polynomial depth, which is certainly not true of degree, our scheme can efficiently evaluate arbitrary circuits without resorting to bootstrapping.

5.5. Security

The security of FHE follows by a standard hybrid argument from the security of E, the basic scheme described in Section 4.1. For the sake of completeness, we will present the full reduction from the hardness of GLWE.

THEOREM 5.7 (SECURITY). *Let $n = n(\lambda)$, $d = d(\lambda)$, $q = q(\lambda)$, $\chi = \chi(\lambda)$, and $L = L(\lambda)$ be functions of the security parameter. Let $N := n \cdot \text{polylog}(q, \lambda)$. The scheme FHE is CPA secure under the GLWE assumption with parameters n, d, q and χ .*

At a high level, the idea is this: the view of a CPA adversary for our scheme is very similar to that for the basic scheme E, except that our adversary also gets to see the key-switching parameters. However, the key-switching parameters are made up of a sequence of outputs of the SwitchKeyGen algorithm which, by Lemma 4.3, are indistinguishable from uniform. A formal proof follows by a hybrid argument.

PROOF. We proceed with the proof using a sequence of hybrids. Let \mathcal{A} be an IND-CPA adversary for FHE. We consider a series of hybrids, where $\text{Adv}_H[\mathcal{A}]$ denotes the success probability of \mathcal{A} in hybrid H .

— *Hybrid H_0 .* This is identical to the IND-CPA game, where the adversary gets a properly distributed public key, generated by FHE.KeyGen, and an encryption of either 0 or 1 computed using FHE.Enc. Recall that the public key consists of the following.

- A matrix $\mathbf{A}_L \leftarrow \text{E.PublicKeyGen}(\mathbf{s}_L)$.
- Pairs $(\mathbf{A}_\ell, \tau_{\mathbf{s}_{(\ell+1)} \rightarrow \mathbf{s}_\ell})$ for $\ell = L - 1$ downto 0, where $\mathbf{A}_\ell \leftarrow \text{E.PublicKeyGen}(\mathbf{s}_\ell)$ and $\tau_{\mathbf{s}_{(\ell+1)} \rightarrow \mathbf{s}_\ell} \leftarrow \text{SwitchKeyGen}(\mathbf{s}_{\ell+1} \otimes \mathbf{s}_{\ell+1}, \mathbf{s}_\ell)$.

For the sake of contradiction, assume that there is a polynomial function $p(\cdot)$ such that

$$\text{Adv}_{H_0}[\mathcal{A}] := |\Pr[\mathcal{A}(pk, \text{FHE.Enc}(pk, \mu_0) = 1] - \Pr[\mathcal{A}(pk, \text{FHE.Enc}(pk, \mu_1) = 1]| > 1/p(\lambda). \quad (3)$$

— *Hybrid H_ℓ , for $\ell \in [L]$.* The hybrids H_ℓ are identical to $H_{\ell-1}$ except that the public key $\mathbf{A}_{\ell-1}$ and the key-switching parameter $\tau_{\mathbf{s}_\ell \rightarrow \mathbf{s}_{\ell-1}}$ are chosen to be uniformly random from the appropriate domains. We claim that

$$|\text{Adv}_{H_\ell}[\mathcal{A}] - \text{Adv}_{H_{\ell-1}}[\mathcal{A}]| \leq \text{negl}(\lambda). \quad (4)$$

Assume for contradiction that there is an adversary \mathcal{A} that distinguishes between H_ℓ and $H_{\ell-1}$ with nonnegligible advantage $1/q(\lambda)$ for some polynomial $q(\cdot)$. Then, we construct a distinguisher \mathcal{B} that distinguishes between a correctly generated pair $(\mathbf{A}_\ell, \tau_{\mathbf{s}_{(\ell+1)} \rightarrow \mathbf{s}_\ell})$ and a uniformly random tuple of the same dimensions, for some $\mathbf{s}_{\ell-1}$. Such a distinguisher cannot exist, by Lemma 4.3.

The distinguisher \mathcal{B} picks vectors $\mathbf{s}_\ell, \dots, \mathbf{s}_{L-1}$ uniformly at random and generates pairs $(\mathbf{A}_\ell, \tau_{\mathbf{s}_{(\ell+1)} \rightarrow \mathbf{s}_\ell})$, where

$$\mathbf{A}_\ell \leftarrow \mathsf{E}.\mathsf{PublicKeyGen}(\mathbf{s}_\ell) \text{ and } \tau_{\mathbf{s}_{(\ell+1)} \rightarrow \mathbf{s}_\ell} \leftarrow \mathsf{SwitchKeyGen}(\mathbf{s}_{\ell+1} \otimes \mathbf{s}_{\ell+1}, \mathbf{s}_\ell),$$

and also generates $\mathbf{A}_L \leftarrow \mathsf{E}.\mathsf{PublicKeyGen}(\mathbf{s}_L)$. \mathcal{B} then sets $(\mathbf{A}_{\ell-1}, \tau_{\mathbf{s}_\ell \rightarrow \mathbf{s}_{\ell-1}})$ to be its own challenge input, and chooses the rest of the components of the public key at random from their respective domains. Finally, it encrypts a random bit b and feeds it to the distinguisher \mathcal{A} . If \mathcal{A} returns a bit $b' = b$, then output 1, else output 0.

It is easy to see that the distribution generated by \mathcal{B} is either $H_{\ell-1}$ or H_ℓ depending on whether its challenge input was correctly generated or random. This proves our claim.

Note that in hybrid H_L , all the pairs $(\mathbf{A}_{\ell-1}, \tau_{\mathbf{s}_\ell \rightarrow \mathbf{s}_{\ell-1}})$ are uniformly random for every $\ell \in [L]$.

—*Hybrid H_{L+1}* . Hybrid H_{L+1} is identical to H_L except that the matrix \mathbf{A}_L is chosen uniformly at random rather than being generated as an output of $\mathsf{E}.\mathsf{PublicKeyGen}(\mathbf{s}_L)$. It follows that

$$|\mathsf{Adv}_{H_{L+1}}[\mathcal{A}] - \mathsf{Adv}_{H_L}[\mathcal{A}]| \leq \mathsf{negl}(\lambda) \quad (5)$$

by an argument essentially identical to the preceding one.

Note that in H_{L+1} , all the elements of the public key are uniformly random and independent of the message. Thus, invoking the semantic security of the basic encryption scheme E , we get that

$$\mathsf{Adv}_{H_{L+1}}[\mathcal{A}] = \mathsf{Adv}_{\mathsf{E}, cpa}[\mathcal{A}] = \mathsf{negl}(\lambda). \quad (6)$$

Putting Equations (4), (5), and (6) together, we get

$$\mathsf{Adv}_{H_0}[\mathcal{A}] \leq \mathsf{Adv}_{H_{L+1}}[\mathcal{A}] + \sum_{\ell=0}^L |\mathsf{Adv}_{H_\ell}[\mathcal{A}] - \mathsf{Adv}_{H_{\ell+1}}[\mathcal{A}]| = \mathsf{negl}(\lambda),$$

contradicting Equation (3). The theorem follows. \square

6. OPTIMIZATIONS: REDUCING THE PER-GATE COMPUTATION OVERHEAD

Despite the fact that the RLWE-based version of our new FHE scheme has per-gate computation only quasilinear in the security parameter, it has a rather large dependence on the number of levels in the circuit (see Theorem 5.6). In this section, we show how to reduce the per-gate computation overhead to quasilinear in the security parameter, independent of the number of levels in the circuit to be evaluated. To this end, we present several significant optimizations.

Our first optimization (Section 6.1) is *batching*. Batching allows us to evaluate a function f homomorphically in parallel on $\ell = \Omega(\lambda)$ blocks of encrypted data, paying only a polylogarithmic (in the security parameter λ) overhead over evaluating f on the unencrypted data. (The overhead is still polynomial in the depth L of the circuit computing f .) Thus, batching allows us to reduce the per-gate computation from quasilinear in the security parameter to polylogarithmic, in an amortized sense. Batching works essentially by packing multiple plaintexts into each ciphertext.

Second, we bring *bootstrapping* back into the picture as an optimization rather than a necessity (Section 6.2). Bootstrapping allows us to achieve per-gate computation quasiquadratic in the security parameter, independent of the number levels in the circuit being evaluated. Although the dependence on the security parameter has worsened from $\tilde{O}(\lambda)$ to $\tilde{O}(\lambda^2)$, the key win here is that the per-gate computation does not depend on the depth of the circuit any more! This optimization applies in a stand-alone setting where the function is evaluated on a single input.

Finally, we obtain our goal of near-linear per-gate computation overhead (in the security parameter λ) by showing that *batching the bootstrapping function* is a powerful combination (Section 6.3). With this optimization, circuits that are wide on average, namely, ones whose average width is at least λ , can be evaluated homomorphically with only $\tilde{O}(\lambda)$ per-gate computation, independent of the number of levels.

6.1. Batching

Suppose we want to evaluate the same function f on ℓ blocks of encrypted data. (Or, similarly, suppose we want to evaluate the same encrypted function f on ℓ blocks of plaintext data.) Can we do this using less than ℓ times the computation needed to evaluate f on one block of data? Can we batch?

For example, consider a keyword search function that returns ‘1’ if the keyword is present in the data and ‘0’ if it is not. The keyword search function is mostly composed of a large number of equality tests that compare the target word w to all of the different subsequences of data; this is followed up by an OR of the equality test results. All of these equality tests involve running the same w -dependent function on different blocks of data. If we could batch these equality tests, it could significantly reduce the computation needed to perform keyword search homomorphically.

If we use bootstrapping as an optimization (see Section 6.2), then obviously we will be running the decryption function homomorphically on multiple blocks of data—namely, the multiple ciphertexts that need to be refreshed. Can we batch the bootstrapping function? If we could, then we might be able to drastically reduce the average per-gate cost of bootstrapping.

Smart and Vercauteren [2011] were the first to rigorously analyze batching in the context of FHE. In particular, they observed that ideal-lattice-based (and RLWE-based) ciphertexts can have many plaintext slots, associated to the factorization of the plaintext space into algebraic ideals.

When we apply batching to our new RLWE-based FHE scheme, the results are pretty amazing. Evaluating f homomorphically on $\ell = \Omega(\lambda)$ blocks of encrypted data requires only polylogarithmically (in terms of the security parameter λ) more computation than evaluating f on the unencrypted data. (The overhead is still polynomial in the depth L of the circuit computing f .) As we will see later, for circuits whose levels have average width at least λ , batching the bootstrapping function (i.e., batching homomorphic evaluation of the decryption function) allows us to reduce the per-gate computation of our bootstrapped scheme from $\tilde{O}(\lambda^2)$ to $\tilde{O}(\lambda)$ (independent of L).

To make the exposition a bit simpler, in our RLWE-based instantiation where $R = \mathbb{Z}[x]/(x^d + 1)$, we will not use R_2 as our plaintext space, but instead use a plaintext space R_p , prime $p = 1 \bmod 2d$, where we have the isomorphism $R_p \cong R_{p_1} \times \dots \times R_{p_d}$ of many plaintext spaces (think Chinese remaindering), so that evaluating a function once over R_p implicitly evaluates the function many times in parallel over the respective smaller plaintext spaces. The p_i 's will be *ideals* in our ring $R = \mathbb{Z}[x]/(x^d + 1)$. (One could still use R_2 as in Smart and Vercauteren [2011], but the number theory there is a bit more involved.)

6.1.1. Some Number Theory. Let us take a very brief tour of algebraic number theory. Suppose p is a prime number satisfying $p \equiv 1 \pmod{2d}$, and let a be a primitive $2d$ th root of unity modulo p . Then, $x^d + 1$ factors completely into linear polynomials modulo p —in particular, $x^d + 1 = \prod_{i=1}^d (x - a_i) \pmod{p}$, where $a_i = a^{2i-1} \pmod{p}$. In some sense, the converse of the preceding statement is also true, and this is the essence of reciprocity—namely, in the ring $R = \mathbb{Z}[x]/(x^d + 1)$, the prime integer p is not actually prime, but rather it splits completely into prime ideals in R , that is, $p = \prod_{i=1}^d \mathfrak{p}_i$. The ideal \mathfrak{p}_i equals $(p, x - a_i)$, namely, the set of all R -elements that can be expressed as $r_1 \cdot p + r_2 \cdot (x - a_i)$ for some $r_1, r_2 \in R$. Each ideal \mathfrak{p}_i has norm p , that is, roughly speaking, a $1/p$ fraction of R -elements are in \mathfrak{p}_i , or, more formally, the p cosets $0 + \mathfrak{p}_i, \dots, (p-1) + \mathfrak{p}_i$ partition R . The latter implies that $R_{\mathfrak{p}_i} := R/\mathfrak{p}_i R$ is isomorphic to \mathbb{Z}_p . The ideals $R_{\mathfrak{p}_i}$ are relatively prime, and so they behave like relatively prime integers. In particular, the Chinese Remainder Theorem applies $R_p \cong R_{\mathfrak{p}_1} \times \dots \times R_{\mathfrak{p}_d}$, and in turn, $R_p \cong (\mathbb{Z}_p)^d$. For any $x \in R$, we will use $[x]_{\mathfrak{p}_i}$ to indicate that value $k \in \{0, \dots, p\}$ for which $x \in k + \mathfrak{p}_i$.

Although the prime ideals $\{\mathfrak{p}_i\}$ are relatively prime, they are close siblings, and it is easy, in some sense, to switch from one to another. One fact that we will use (when we finally apply batching to bootstrapping) is that, for any i, j , there is an automorphism $\sigma_{i \rightarrow j}$ over R that maps elements of \mathfrak{p}_i to elements of \mathfrak{p}_j . Specifically, $\sigma_{i \rightarrow j}$ works by mapping an R -element $r = r(x) = r_{d-1}x^{d-1} + \dots + r_1x + r_0$ to $r(x^{e_{ij}}) = r_{d-1}x^{e_{ij}(d-1)} \pmod{2d} + \dots + r_1x^{e_{ij}} + r_0 \pmod{(x^d + 1)}$, where e_{ij} is some number relative prime to $2d$. Notice that this automorphism just permutes the coefficients of r and fixes the free coefficient. Notationally, we will use $\sigma_{i \rightarrow j}(\mathbf{v})$ to refer to the vector that results from applying $\sigma_{i \rightarrow j}$ coefficient-wise to \mathbf{v} . Note that $\sigma_{i \rightarrow j}(\cdot)$ permutes the elements of its input vector and possibly flips the sign for some of them.

6.1.2. How Batching Works. We will show that the following holds.

THEOREM 6.1. *Let $p \equiv 1 \pmod{2d}$ be a prime of size polynomial in λ . The RLWE-based instantiation of FHE using the ring $R = \mathbb{Z}[x]/(x^d + 1)$ can be adapted to use the plaintext space $R_p = \otimes_{i=1}^d R_{\mathfrak{p}_i}$ while preserving correctness and the same asymptotic performance. For any boolean circuit f of depth L , the scheme can homomorphically evaluate f on ℓ sets of inputs with per-gate computation $\tilde{O}(\lambda \cdot L^3 / \min\{d, \ell\})$.*

When $\ell = \Omega(\lambda)$, and taking $d = \tilde{\Omega}(\lambda)$, the per-gate computation is only polylogarithmic in the security parameter (still cubic in L).

PROOF. Deploying batching inside our scheme FHE is quite straightforward. First, we pick a prime $p \equiv 1 \pmod{2d}$ of size polynomial in the security parameter. (One should exist under the GRH.)

The next step is simply to recognize that our scheme FHE works just fine when we replace the original plaintext space R_2 with R_p . There is nothing especially magical about the number 2. In the basic scheme E described in Section 4.1, $E.\text{PublicKeyGen}(\mathsf{params}, \mathsf{sk})$ is modified in the obvious way so that $\mathbf{A} \cdot \mathbf{s} = p \cdot \mathbf{e}$, rather than $2 \cdot \mathbf{e}$. (This modification induces a similar modification in SwitchKeyGen .) Decryption becomes $m = [\lfloor \langle \mathbf{c}, \mathbf{s} \rangle \rfloor_q]_p$. Homomorphic operations use mod- p gates rather than boolean gates, and it is easy (if desired) to emulate boolean gates with mod- p gates (e.g., we can compute $\text{XOR}(a, b)$ for $a, b \in \{0, 1\}^2$ using mod- p gates for any p as $a + b - 2ab$). For modulus switching, we use $\text{Scale}(\mathbf{c}_1, q_j, q_{j-1}, p)$ rather than $\text{Scale}(\mathbf{c}_1, q_j, q_{j-1}, 2)$. The larger rounding error from this new scaling procedure increases the noise slightly, but this additive noise is still polynomial in the security parameter and the number of levels, and thus is still consistent with our setting of parameters. In short, FHE can easily be adapted to work with a plaintext space R_p for p of polynomial size.

The final step is simply to recognize that, by the Chinese Remainder Theorem, evaluating an arithmetic circuit over R_p on input $\mathbf{x} \in R_p^n$ implicitly evaluates, for each i , the same arithmetic circuit over R_{p_i} on input \mathbf{x} projected down to $R_{p_i}^n$. The evaluations modulo the various prime ideals do not “mix” or interact with each other. \square

6.2. Bootstrapping as an Optimization

Bootstrapping is no longer strictly necessary to achieve leveled FHE. However, in some settings, it may have some advantages.

- *Performance.* The per-gate computation is independent of the depth of the circuit being evaluated.
- *Flexibility.* Assuming circular security, a bootstrapped scheme can perform homomorphic evaluations indefinitely without needing to specify in advance, during Setup, a bound on the number of circuit levels.
- *Memory.* Bootstrapping permits short ciphertexts (e.g., encrypted using AES other space-efficient cryptosystem) to be decompressed to longer ciphertexts that permit homomorphic operations. Bootstrapping thus allows us to save memory by storing data encrypted in the compressed form, while retaining the ability to perform homomorphic operations.

Here, we revisit bootstrapping, viewing it as an optimization rather than a necessity. We also reconsider the scheme FHE that we described in Section 4, viewing the scheme not as an end in itself, but rather as a very powerful SWHE whose performance degrades polynomially in the depth of the circuit being evaluated, as opposed to previous SWHE schemes whose performance degrades polynomially in the degree. In particular, we analyze how efficiently it can evaluate its decryption function as needed to bootstrap. Not surprisingly, our faster SWHE scheme can also bootstrap faster. The decryption function has only logarithmic depth and can be evaluated homomorphically in time quasiquadratic in the security parameter (for the RLWE instantiation), giving a bootstrapped scheme with quasiquadratic per-gate computation overall.

To apply the Bootstrapping Theorem, we need to bound the circuit complexity of the decryption of our scheme. We start by recalling that the decryption function is $m = [\lfloor \langle \mathbf{c}, \mathbf{s} \rangle \rfloor_q]_2$. Suppose that we are given the bits (elements in R_2) of \mathbf{s} as input, and we want to compute $[\lfloor \langle \mathbf{c}, \mathbf{s} \rangle \rfloor_q]_2$ using an arithmetic circuit that has Add and Mult gates over R_2 . (When we bootstrap, of course we are given the bits of \mathbf{s} in encrypted form.) Note that we will run the decryption function homomorphically on level-0 ciphertexts, that is, when q is small, only polynomial in the security parameter. What is the complexity of this circuit? Most importantly for our purposes, what is its depth and size? The answer is that we can perform decryption with $\tilde{O}(\lambda)$ computation and $O(\log \lambda)$ depth. Thus, in the RLWE instantiation, we can evaluate the decryption function homomorphically using our new scheme with quasiquadratic computation. (For the LWE instantiation, the bootstrapping computation is quasiquadratic.)

LEMMA 6.2. *Let $n = O(\lambda)$, $q = \text{poly}(\lambda)$ and let $\mathbf{c} \in \mathbb{Z}_q^n$. Consider the function $D_{\mathbf{c}}(\mathbf{s}) = [\lfloor \langle \mathbf{c}, \mathbf{s} \rangle \rfloor_q]_2$, for inputs $\mathbf{s} \in \mathbb{Z}_q^n$, where the elements of \mathbf{s} are represented using standard binary representation. Then for all \mathbf{c} , $D_{\mathbf{c}}(\cdot)$ can be computed by a circuit of size $\tilde{O}(\lambda)$ and depth $O(\log \lambda)$.*

PROOF. Obviously, each product $\mathbf{c}[i] \cdot \mathbf{s}[i]$ can be written as the sum of at most $\log q$ shifts of $\mathbf{s}[i]$. These horizontal shifts of $\mathbf{s}[i]$ use at most $2 \log q$ columns. Thus, $\langle \mathbf{c}, \mathbf{s} \rangle$ can be written as the sum of $n \cdot \log q$ numbers, where each number has $2 \log q$ digits. As discussed in Gentry [2009b], we can use the three-for-two trick, which takes as input three numbers in binary (of arbitrary length) and outputs (using constant depth) two binary

numbers with the same sum. Thus, with $O(\log(n \cdot \log q)) = O(\log n + \log \log q)$ depth and $O(n \log^2 q)$ computation, we obtain two numbers with the desired sum, each having $O(\log n + \log q)$ bits. We can sum the final two numbers with $O(\log \log n + \log \log q)$ depth and $O(\log n + \log q)$ computation. So far, we have used depth $O(\log n + \log \log q)$ and $O(n \log^2 q)$ computation to compute $\langle \mathbf{c}, \mathbf{s} \rangle$. Reducing this value modulo q is an operation akin to division for which there are circuits of size $\text{polylog}(q)$ and depth $\log \log q$. Finally, reducing modulo 2 just involves dropping the most significant bits. Overall, since we are interested only in the case where $\log q = O(\log \lambda)$, we have that decryption requires $\tilde{O}(\lambda)$ computation and depth $O(\log \lambda)$. \square

The preceding lemma generalizes to GLWE as follows.

LEMMA 6.3. *Let $q = \text{poly}(\lambda)$, n, d such that d is a power of 2 and $n \cdot d = O(\lambda)$, $R = \mathbb{Z}[x]/(x^d + 1)$, and $\mathbf{c} \in R_q^n$. Consider the function $D_{\mathbf{c}}(\mathbf{s}) = [\langle \mathbf{c}, \mathbf{s} \rangle]_q$, for inputs $\mathbf{s} \in R_q^n$, where the coefficients of \mathbf{s} are represented using standard binary representation. Then for all \mathbf{c} , $D_{\mathbf{c}}(\cdot)$ can be computed by a circuit of size $\tilde{O}(\lambda)$ and depth $O(\log \lambda)$.*

PROOF. The proof is practically identical to Lemma 6.2, except we use DFT to multiply elements in R instead of “standard” integer multiplication. Since all roots of $x^d + 1$ are $2d$ th roots of unity, we can use FFT to achieve multiplication in logarithmic depth and quasilinear size (in $d \log q$).

The procedure thus proceeds exactly as in Lemma 6.2, except the last operations of taking modulo q and then modulo 2 are performed d times, but since these only require $\log(q)$ depth and $\text{polylog}(q)$ size, the bound still holds. \square

We can now apply the Bootstrapping Theorem to our scheme from Section 4. The relevant values of n, d, q when applying Lemma 6.3 are those of the last level (level L) of the scheme, since these are the values that are actually being decrypted. While modulus reduction guarantees that the values of q_L, n_L are independent of L , it is normally the case that d_L depends on L (think, e.g., about the case of $n = 1$). It may seem therefore that the decryption depth depends on L contrary to the compactness requirement. Using the techniques of Section 3, we can convert d_L back down to d_0 , which is independent of L . (We remark, however, that even if we don’t make the conversion, the dependence on L is small enough that our scheme is bootstrappable.)

Putting the pieces together, we get the following.

THEOREM 6.4. *The scheme FHE with parameters n, d, L such that $n \cdot d = O(\lambda)$ and $L = O(\log \lambda)$ is bootstrappable. The per-gate computation is $\tilde{O}(\lambda^4)$ for the LWE case ($d = 1$) and $\tilde{O}(\lambda^2)$ for the RLWE case ($n = 1$).*

PROOF. The theorem is a straightforward derivation from combining the bootstrapping theorem, Lemma 6.3, and Theorem 5.6. \square

6.3. Batching the Bootstrapping Operation

Suppose that we are evaluating a circuit homomorphically, that we are currently at a level in the circuit that has at least d gates (where d is the dimension of our ring), and that we want to bootstrap (refresh) all of the ciphertexts corresponding to the respective wires at that level. That is, we want to homomorphically evaluate the decryption function at least d times in parallel. This seems like an ideal place to apply batching.

However, there are some nontrivial problems. In Section 6.1, our focus was rather limited. For example, we did not consider whether homomorphic operations could continue after the batched computation. Indeed, at first glance, it would appear that

homomorphic operations cannot continue, since, after batching, the encrypted data is partitioned into non-interacting, relatively-prime plaintext slots, whereas the whole point of homomorphic encryption is that the encrypted data can interact (within a common plaintext slot). Similarly, we did not consider homomorphic operations before the batched computation. Somehow, we need the input to the batched computation to come pre-partitioned into the different plaintext slots.

What we need are Pack and Unpack functions that allow the batching procedure to interface with normal homomorphic operations. One may think of the Pack and Unpack functions as an on-ramp to and an exit-ramp from the “fast lane” of batching. Let us say that normal homomorphic operations will always use the plaintext slot $R_{\mathfrak{p}_1}$. Roughly, the Pack function should take a bunch of ciphertexts $\mathbf{c}_1, \dots, \mathbf{c}_d$ that encrypt messages $m_1, \dots, m_d \in \mathbb{Z}_p$ under key \mathbf{s}_1 for modulus q and plaintext slot $R_{\mathfrak{p}_1}$, and then aggregate them into a single ciphertext \mathbf{c} under some possibly different key \mathbf{s}_2 for modulus q , so that correctness holds with respect to all of the different plaintext slots, that is, $m_i = [\langle \mathbf{c}, \mathbf{s}_2 \rangle]_q|_{\mathfrak{p}_i}$ for all i . The Pack function thus allows normal homomorphic operations to feed into the batch operation. The Unpack function should accept the output of a batched computation, namely, a ciphertext \mathbf{c}' such that $m_i = [\langle \mathbf{c}', \mathbf{s}'_1 \rangle]_q|_{\mathfrak{p}_i}$ for all i , and then de-aggregate this ciphertext by outputting ciphertexts $\mathbf{c}'_1, \dots, \mathbf{c}'_d$ under some possibly different common secret key \mathbf{s}'_2 such that $m_i = [\langle \mathbf{c}', \mathbf{s}'_2 \rangle]_q|_{\mathfrak{p}_1}$ for all i . Now that all of the ciphertexts are under a common key and plaintext slot, normal homomorphic operations can resume. With such Pack and Unpack functions, we could indeed batch the bootstrapping operation. For circuits of large width (say, at least d), we could reduce the per-gate bootstrapping computation by a factor of d , making it only quasilinear in λ . Assuming the Pack and Unpack functions have complexity at most quasiquadratic in d (per-gate this is only quasilinear, since Pack and Unpack operate on d gates), the overall per-gate computation of a batched-bootstrapped scheme becomes only quasilinear.

Here, we describe suitable Pack and Unpack functions. These functions will make heavy use of the automorphisms $\sigma_{i \rightarrow j}$ over R that map elements of \mathfrak{p}_i to elements of \mathfrak{p}_j . (See Section 6.1.1.) We note that Smart and Vercauteren [2011] used these automorphisms to construct something similar to our Pack function (though for unpacking they resorted to bootstrapping). We also note that Lyubashevsky et al. [2010] used these automorphisms to permute the ideal factors q_i of the modulus q , which was an essential tool toward their proof of the pseudorandomness of RLWE.

Toward Pack and Unpack procedures, we begin with the observation that if m is encoded as a number in $\{0, \dots, p - 1\}$ and if $m = [\langle \mathbf{c}, \mathbf{s} \rangle]_q|_{\mathfrak{p}_i}$, then $m = [\langle \sigma_{i \rightarrow j}(\mathbf{c}), \sigma_{i \rightarrow j}(\mathbf{s}) \rangle]_q|_{\mathfrak{p}_j}$. That is, we can switch the plaintext slot but leave the decrypted message unchanged by applying the same automorphism to the ciphertext and the secret key. (These facts follow from the fact that $\sigma_{i \rightarrow j}$ is a homomorphism, that it maps elements of \mathfrak{p}_i to elements of \mathfrak{p}_j , and that it fixes integers.) Of course, then we have a problem: the ciphertext is now under a different key, whereas we may want the ciphertext to be under the same key as other ciphertexts. To get the ciphertexts to be back under the same key, we simply use the SwitchKey algorithm to switch all of the ciphertexts to a new common key.

Some technical remarks before we describe Pack / Unpack more formally: We mention again that E.PublicKeyGen is modified in the obvious way so that $\mathbf{A} \cdot \mathbf{s} = p \cdot \mathbf{e}$ rather than $2 \cdot \mathbf{e}$, and that this modification induces a similar modification in SwitchKeyGen. Also, let $u \in R$ be a short element such that $u \in 1 + \mathfrak{p}_1$ and $u \in \mathfrak{p}_j$ for all $j \neq 1$. It is obvious that such a u with coefficients in $(-p/2, p/2]$ can be computed efficiently by first picking any element u' such that $u' \in 1 + \mathfrak{p}_1$ and $u' \in \mathfrak{p}_j$ for all $j \neq 1$, and then reducing the coefficients of u' modulo p .

PackSetup($\mathbf{s}_1, \mathbf{s}_2$). Takes as input two secret keys $\mathbf{s}_1, \mathbf{s}_2$. For all $i \in [1, d]$, it runs $\tau_{(\sigma_{1 \rightarrow i}(\mathbf{s}_1)) \rightarrow \mathbf{s}_2} \leftarrow \text{SwitchKeyGen}(\sigma_{1 \rightarrow i}(\mathbf{s}_1), \mathbf{s}_2)$.

Pack($\{\mathbf{c}_i\}_{i=1}^d, \{\tau_{\sigma_{1 \rightarrow i}(\mathbf{s}_1) \rightarrow \mathbf{s}_2}\}_{i=1}^d$). Takes as input ciphertexts $\mathbf{c}_1, \dots, \mathbf{c}_d$ such that $m_i = [\langle \mathbf{c}_i, \mathbf{s}_1 \rangle]_q$ and $0 = [\langle \mathbf{c}_j, \mathbf{s}_1 \rangle]_q$ for all $j \neq i$, and also some auxiliary information output by PackSetup. For all i , it does the following.

- Computes $\mathbf{c}_i^* \leftarrow \sigma_{1 \rightarrow i}(\mathbf{c}_i)$. (Observe: We have $m_i = [\langle \mathbf{c}_i^*, \sigma_{1 \rightarrow i}(\mathbf{s}_1) \rangle]_q$, while $0 = [\langle \mathbf{c}_j^*, \sigma_{1 \rightarrow i}(\mathbf{s}_1) \rangle]_q$ for all $j \neq i$)
- Runs $\mathbf{c}_i^\dagger \leftarrow \text{SwitchKey}(\tau_{(\sigma_{1 \rightarrow i}(\mathbf{s}_1)) \rightarrow \mathbf{s}_2}, \mathbf{c}_i^*)$. (Observe: Assuming the noise does not wrap, we have that $m_i = [\langle \mathbf{c}_i^\dagger, \mathbf{s}_2 \rangle]_q$ and $0 = [\langle \mathbf{c}_j^\dagger, \mathbf{s}_2 \rangle]_q$ for all $j \neq i$.)

Finally, it outputs $\mathbf{c} \leftarrow \sum_{i=1}^d \mathbf{c}_i^\dagger$. (Observe: Assuming the noise does not wrap, we have $m_i = [\langle \mathbf{c}, \mathbf{s}_2 \rangle]_q$ for all i .)

UnpackSetup($\mathbf{s}_1, \mathbf{s}_2$). Takes as input secret keys $\mathbf{s}_1, \mathbf{s}_2$. For all $i \in [1, d]$, it runs

$$\tau_{\sigma_{i \rightarrow 1}(\mathbf{s}_1) \rightarrow \mathbf{s}_2} \leftarrow \text{SwitchKeyGen}(\sigma_{i \rightarrow 1}(\mathbf{s}_1), \mathbf{s}_2).$$

Unpack($\mathbf{c}, \{\tau_{\sigma_{i \rightarrow 1}(\mathbf{s}_1) \rightarrow \mathbf{s}_2}\}_{i=1}^d$). Takes as input a ciphertext \mathbf{c} such that $m_i = [\langle \mathbf{c}, \mathbf{s}_1 \rangle]_q$ for all i , and also some auxiliary information output by UnpackSetup. For all i , it does the following.

- Computes $\mathbf{c}_i \leftarrow u \cdot \sigma_{i \rightarrow 1}(\mathbf{c})$. (Observe: Assuming the noise does not wrap, $m_i = [\langle \mathbf{c}_i, \sigma_{i \rightarrow 1}(\mathbf{s}_1) \rangle]_q$ and $0 = [\langle \mathbf{c}_j, \sigma_{i \rightarrow 1}(\mathbf{s}_1) \rangle]_q$ for all $j \neq i$)
- Outputs $\mathbf{c}_i^* \leftarrow \text{SwitchKey}(\tau_{\sigma_{i \rightarrow 1}(\mathbf{s}_1) \rightarrow \mathbf{s}_2}, \mathbf{c}_i)$. (Observe: Assuming the noise does not wrap, $m_i = [\langle \mathbf{c}_i^*, \mathbf{s}_2 \rangle]_q$ and $0 = [\langle \mathbf{c}_j^*, \mathbf{s}_2 \rangle]_q$ for all $j \neq i$.)

The properties of the resulting scheme are summarized in the following theorem, which is only stated for the RLWE setting (since batched bootstrapping only works in the ring setting).

THEOREM 6.5. *The scheme FHE with parameters n, d, L such that $n = 1$, $d = O(\lambda)$, and $L = O(\log \lambda)$ with batched bootstrapping has per-gate computation $\tilde{O}(\lambda)$ when evaluating circuits of average width $\Omega(\lambda)$.*

PROOF. Consider an evaluation of depth t circuit, and let w_1, \dots, w_t be the width of each of the levels of the circuit. Then, let $w = (1/t) \cdot \sum_{i \in [t]} w_i$ be the average width of the circuit and by the theorem statement $w = \Omega(\lambda)$. The total number of gates is (naturally) $t \cdot w$.

The evaluation of level i of the circuit involves a bootstrapping operation of the $2 \cdot w_i$ input wires into the gates, in addition to w_i gate evaluations.

We pack the $2 \cdot w_i$ standalone ciphertexts into $\lceil 2w_i/d \rceil \leq 2w_i/d + 1$ packed ciphertexts, where $d = \Omega(\lambda)$ is the parameter of the ring that also determines the number of ciphertexts that can be packed. The cost of this packing/unpacking is linear in the input length which is $2w_i \cdot \tilde{O}(\lambda)$.

For each packed ciphertext, we perform bootstrapping which involves a homomorphic evaluation of the decryption circuit. The total cost per packed ciphertext is thus $\tilde{O}(\lambda^2)$. The total complexity for all packed ciphertexts is therefore at most $(2w_i/d + 1) \cdot \tilde{O}(\lambda^2)$.

Finally, the evaluation of the w_i actual gates has (total) complexity $w_i \cdot \tilde{O}(\lambda)$.

Summing all of this, we get that the total complexity of evaluating level i of the circuit is

$$2w_i \cdot \tilde{O}(\lambda) + (2w_i/d + 1) \cdot \tilde{O}(\lambda^2) + w_i \cdot \tilde{O}(\lambda) = (w_i + \lambda + w_i\lambda/d) \cdot \tilde{O}(\lambda).$$

Summing over all t levels, we get that the total complexity of evaluating the entire circuit is at most

$$(t \cdot w + t \cdot \lambda + t \cdot w\lambda/d) \cdot \tilde{O}(\lambda),$$

and the per-gate cost is obtained by dividing by $t \cdot w$, and recalling that $w, d = \Omega(\lambda)$:

$$(1 + \lambda/w + \lambda/d) \cdot \tilde{O}(\lambda) = \tilde{O}(\lambda).$$

The theorem thus follows. \square

ACKNOWLEDGMENTS

We thank Carlos Aguilar Melchor, Boaz Barak, Shafi Goldwasser, Shai Halevi, Chris Peikert, Nigel Smart, and Jiang Zhang for their helpful discussions and insights.

REFERENCES

- Miklós Ajtai, Ravi Kumar, and D. Sivakumar. 2001. A sieve algorithm for the shortest lattice vector problem. In *Proceedings of the 33rd Annual ACM Symposium on Theory of Computing (STOC'01)*. ACM, 601–610.
- Benny Applebaum, David Cash, Chris Peikert, and Amit Sahai. 2009. Fast cryptographic primitives and circular-secure encryption based on hard learning problems. In *Proceedings of the 29th Annual International Cryptology Conference (CRYPTO'09)*. Lecture Notes in Computer Science, vol. 5677, Springer, Berlin, 595–618.
- Dan Boneh, Eu-Jin Goh, and Kobbi Nissim. 2005. Evaluating 2-DNF formulas on ciphertexts. In *Proceedings of the 2nd Theory of Cryptography Conference (TCC'05)*. Lecture Notes in Computer Science, vol. 3378, Springer, Berlin, 325–342.
- Zvika Brakerski. 2012. Fully homomorphic encryption without modulus switching from classical GapSVP. In *Proceedings of the 32nd Annual Cryptology Conference (CRYPTO'12)*. Lecture Notes in Computer Science, vol. 7417, Springer, Berlin, 868–886.
- Zvika Brakerski and Vinod Vaikuntanathan. 2011a. Fully homomorphic encryption from ring-LWE and security for key dependent messages. In *Proceedings of the 31st Annual Cryptology Conference (CRYPTO'11)*. Lecture Notes in Computer Science, vol. 6841, Springer, Berlin, 505–524.
- Zvika Brakerski and Vinod Vaikuntanathan. 2011b. Efficient fully homomorphic encryption from (standard) LWE. In *Proceedings of the IEEE 52nd Annual Symposium on Foundations of Computer Science (FOCS'11)*. IEEE, 97–106.
- Jean-Sébastien Coron, Avradip Mandal, David Naccache, and Mehdi Tibouchi. 2011. Fully homomorphic encryption over the integers with shorter public keys. In *Proceedings of the 31st Annual Cryptology Conference (CRYPTO'11)*, Phillip Rogaway Ed., Lecture Notes in Computer Science, vol. 6841, Springer, Berlin, 487–504.
- Craig Gentry. 2009a. A fully homomorphice encryption scheme. Ph.D. dissertation, Stanford University, Stanford, CA. crypto.stanford.edu/craig/.
- Craig Gentry. 2009b. Fully homomorphic encryption using ideal lattices. In *Proceedings of the 41st Annual ACM Symposium on Theory of Computing (STOC'09)*. Michael Mitzenmacher Ed., ACM, 169–178.
- Craig Gentry and Shai Halevi. 2011a. Fully homomorphic encryption without squashing using depth-3 arithmetic circuits. In *Proceedings of the IEEE 52nd Annual Symposium on Foundations of Computer Science (FOCS'11)*. Rafail Ostrovsky Ed., IEEE, 107–109.
- Craig Gentry and Shai Halevi. 2011b. Implementing gentry's fully-homomorphic encryption scheme. In *Proceedings of the 30th Annual International Conference on the Theory and Applications of Cryptographic Techniques (EUROCRYPT'11)*. Lecture Notes in Computer Science, vol. 6632, Springer, Berlin, 29–148.
- Craig Gentry, Shai Halevi, Chris Peikert, and Nigel P. Smart. 2012a. Ring switching in BGV-style homomorphic encryption. In *Proceedings of the 8th International Conference on Security and Cryptography for Networks (SCN'12)*. Ivan Visconti and Roberto De Prisco Eds., Lecture Notes in Computer Science, vol. 7485, Springer, Berlin, 19–37.
- Craig Gentry, Shai Halevi, and Nigel P. Smart. 2012b. Better bootstrapping in fully homomorphic encryption. In *Proceedings of the 15th International Conference on Practice and Theory in Public Key Cryptography (PKC'12)*. Springer, Berlin, 43–60.

- Cryptography (PKC'12)*. Marc Fischlin, Johannes Buchmann, and Mark Manulis Eds., Lecture Notes in Computer Science, vol. 7293, Springer, Berlin, 1–16.
- Craig Gentry, Shai Halevi, and Nigel P. Smart. 2012c. Fully homomorphic encryption with polylog overhead. In *Proceedings of the 31st Annual International Conference on the Theory and Applications of Cryptographic Techniques (EUROCRYPT'12)*. David Pointcheval and Thomas Johansson Eds., Lecture Notes in Computer Science, vol. 7237, Springer, Berlin, 465–482.
- Craig Gentry, Shai Halevi, and Nigel P. Smart. 2012d. Homomorphic evaluation of the AES circuit. In *Proceedings of the 32nd Annual Cryptology Conference (CRYPTO'12)*. Reihaneh Safavi-Naini and Ran Canetti Eds., Lecture Notes in Computer Science, vol. 7417, Springer, Berlin, 850–867.
- Craig Gentry, Shai Halevi, and Vinod Vaikuntanathan. 2010. i -Hop homomorphic encryption and rerandomizable Yao circuits. In *Proceedings of the 30th Annual Cryptology Conference (CRYPTO'10)*. Lecture Notes in Computer Science, vol. 6223, Springer, Berlin, 155–172.
- Shafi Goldwasser and Silvio Micali. 1984. Probabilistic encryption. *J. Comput. Syst. Sci.* 28, 2, 270–299.
- Shai Halevi and Victor Shoup. 2013. HElib: An implementation of homomorphic encryption. <https://github.com/shaih/HElib>.
- Yuval Ishai and Anat Paskin. 2007. Evaluating branching programs on encrypted data. In *Proceedings of the 4th Theory of Cryptography Conference (TCC'07)*. Salil P. Vadhan Ed., Lecture Notes in Computer Science, vol. 4392, Springer, Berlin, 575–594.
- Kristin Lauter, Michael Naehrig, and Vinod Vaikuntanathan. 2011. Can homomorphic encryption be practical? In *Proceedings of the 3rd ACM Cloud Computing Security Workshop (CCSW'11)*. ACM, 113–124.
- Vadim Lyubashevsky, Chris Peikert, and Oded Regev. 2010. On ideal lattices and learning with errors over rings. In *Proceedings of the 29th Annual International Conference on the Theory and Applications of Cryptographic Techniques (EUROCRYPT'10)*. Lecture Notes in Computer Science, vol. 6110, Springer, Berlin, 1–23.
- Carlos Aguilera Melchor, Philippe Gaborit, and Javier Herranz. 2010. Additively homomorphic encryption with d -operand multiplications. In *Proceedings of the 30th Annual Cryptology Conference (CRYPTO'10)*. Tal Rabin Ed., Lecture Notes in Computer Science, vol. 6223, Springer, Berlin, 138–154.
- Daniele Micciancio. 2007. Generalized compact knapsacks, cyclic lattices, and efficient one-way functions. *Computat. Complex.* 16, 4 (2007), 365–411.
- Daniele Micciancio and Panagiotis Voulgaris. 2010. A deterministic single exponential time algorithm for most lattice problems based on voronoi cell computations. In *Proceedings of the 42nd Annual ACM Symposium on Theory of Computing (STOC'10)*. Leonard J. Schulman Ed., ACM, 351–358.
- Chris Peikert. 2009. Public-key cryptosystems from the worst-case shortest vector problem: Extended abstract. In *Proceedings of the 41st Annual ACM Symposium on Theory of Computing (STOC'09)*. ACM, 333–342.
- Oded Regev. 2005. On lattices, learning with errors, random linear codes, and cryptography. In *Proceedings of the 37th Annual ACM Symposium on Theory of Computing (STOC'05)*. Harold N. Gabow and Ronald Fagin Eds., ACM, 84–93.
- Oded Regev. 2010. Oded Regev. 2010. The learning with errors problem (invited survey). In *Proceedings of the IEEE Conference on Computational Complexity (CCC'10)*. IEEE Computer Society, 191–204.
- Ron Rivest, Leonard Adleman, and Michael L. Dertouzos. 1978. On data banks and privacy homomorphisms. In *Foundations of Secure Computation*. Academic Press, Inc., Orlando, FL, 169–180.
- Claus-Peter Schnorr. 1987. A hierarchy of polynomial time lattice basis reduction algorithms. *Theor. Comput. Sci.* 53, 201–224.
- Nigel P. Smart and Frederik Vercauteren. 2010. Fully homomorphic encryption with relatively small key and ciphertext sizes. In *Proceedings of the 13th International Conference on Practice and Theory in Public Key Cryptography (PKC'10)*. Lecture Notes in Computer Science, vol. 6056, Springer, Berlin, 420–443.
- Nigel P. Smart and Frederik Vercauteren. 2011. Fully homomorphic SIMD operations. *IACR Cryptol. ePrint Archive* 2011, 133.
- Damien Stehlé and Ron Steinfeld. 2010. Faster fully homomorphic encryption. In *Proceedings of the 16th International Conference on the Theory and Application of Cryptology and Information Security (ASIACRYPT'10)*. Lecture Notes of Computer Science, vol. 6477, Springer, Berlin, 377–394.
- Marten van Dijk, Craig Gentry, Shai Halevi, and Vinod Vaikuntanathan. 2010. Fully homomorphic encryption over the integers. In *Proceedings of the 29th Annual International Conference on the Theory and Applications of Cryptographic Techniques (EUROCRYPT'10)*. Lecture Notes in Computer Science, vol. 6110, Springer, Berlin, 24–43.

Received January 2013; revised March 2013; accepted April 2014